

MorphoST: A Simple $E(2)$ -Invariant Transformer for Generalizable Gene Expression Prediction from Histology

Anonymous WACV submission

Paper ID ****

Abstract

001 Predicting spatially resolved gene expression from
002 hematoxylin-and-eosin (H&E) histology could reduce
003 reliance on costly spatial transcriptomics assays. Re-
004 cent methods reach strong accuracy with whole-slide
005 generative modeling, count-distribution priors, and
006 frame-averaged equivariant encoders, but these com-
007 ponents are typically trained and evaluated within a
008 single tissue cohort and add substantial implemen-
009 tation and inference complexity. A more demand-
010 ing requirement is generalization across samples: pre-
011 dicting expression for slides from unseen organs, pa-
012 tients, and acquisition batches. We ask whether a de-
013 liberately simple discriminative model can meet both
014 goals. We introduce MorphoST, a direct-regression
015 transformer that predicts expression at all spots of
016 a slide in a single forward pass. Its spatial pathway
017 uses coordinates only through pairwise distances, mak-
018 ing it invariant to rotations, reflections, and transla-
019 tions of the coordinate system ($E(2)$) for fixed image
020 features; each layer fuses per-spot morphology with
021 distance-biased local attention, global self-attention,
022 and an attention-pooled slide representation. Un-
023 der matched frozen UNI features, splits, and gene
024 panels, we evaluate MorphoST in two complemen-
025 tary regimes—intra-cohort (the standard per-cohort
026 protocol) and the harder cross-organ transfer (leave-
027 one-organ-out and pooled)—against direct, retrieval-
028 based, and flow-matching baselines on HEST-1k and
029 STImage-1K4M, reporting Pearson, Spearman, and
030 MSE. We further analyze learned representations, re-
031 cover spatial domains from predictions, visualize pre-
032 dicted expression maps, and measure inference scaling.
033 These experiments test whether direct spatial regres-
034 sion is a simple, reproducible, and generalizable base-
035 line for histology-based gene expression prediction.

1. Introduction

Spatial transcriptomics (ST) measures gene expression while retaining tissue organization, enabling the study of spatial heterogeneity, the tumor microenvironment, and cell-cell interactions [10]. However, ST assays remain costly and experimentally demanding. This has motivated methods that infer spatial gene expression from routinely acquired hematoxylin-and-eosin (H&E) whole-slide images (WSIs). Recent approaches combine frozen pathology foundation-model features [1] with spatial prediction architectures. STFlow [4], for example, formulates whole-slide prediction using flow matching, a zero-inflated negative-binomial (ZINB) prior, and frame averaging for $SE(2)$ equivariance.

These components provide useful inductive biases but also increase implementation and inference complexity, and they are usually trained separately for each tissue cohort. Yet the setting that matters clinically is cross-sample generalization: a model trained on available tissue should transfer to slides from new organs, patients, and batches (Fig. 1). Prior work shows that models which are strong within a cohort can degrade sharply when transferred to unseen slices, because tissue composition and staining vary across samples. This motivates a focused question: can a direct-regression model, supplied with appropriate local and global spatial context, match generative approaches within a cohort and remain robust across organs? A positive result would provide a complementary baseline that is easier to reproduce, analyze, and extend.

We present MorphoST (Morphology-aware Spatial Transformer). Coordinates enter the network only through pairwise distances, so for fixed image features the prediction is invariant to the 2D Euclidean group $E(2)$ —rotations, reflections, and translations of the coordinate system. Each layer fuses per-spot morphology with three cheap sources of spatial context—distance-biased local attention, global self-attention, and an attention-pooled slide token—and a linear head pre-




Figure 1. Two evaluation regimes. Standard benchmarks train and test within one tissue cohort (left). We additionally test the harder cross-organ transfer (right), where a model must predict expression for slides from organs unseen during training. MorphoST is evaluated in both without per-cohort redesign.

dicts all genes at all spots in one pass, with no generative sampling or count-distribution prior.

Contributions. First, we introduce a single-pass spatial transformer that fuses local, global, and slide-level context through an $E(2)$ -invariant coordinate pathway, giving a simple and reproducible alternative to generative whole-slide prediction (Fig. 2). Second, we evaluate under a unified feature-matched protocol in two regimes—*intra-cohort* and *cross-organ transfer*—on two benchmarks (HEST-1k and STImage-1K4M) and with three metrics (Pearson, Spearman, MSE), against feature-matched baselines. Third, we provide controlled analyses of the design: a factorial context ablation, a study of the learned representation and its ability to recover spatial domains from predicted expression, qualitative expression maps, and measured inference scaling.

2. Related Work

Histology-to-expression prediction. Early spot-based methods regress each spot independently: ST-Net [3] fine-tunes a CNN per patch, and foundation-model probes [1] pair frozen features with a shallow head. These ignore spatial context. Slide-aware methods add it in different ways: HisToGene [7] applies a ViT over all spots with learned positional embeddings; Hist2ST [13] combines a transformer with a graph network over a spatial k NN graph; and TRIPLEX [2] fuses spot, neighbour, and global resolutions with a three-branch attention module. STFlow [4] models the joint slide-level expression distribution using flow matching, and the concurrent MCToGene [11] models high-order many-to-many cross-cell interactions, reporting the current best per-cohort accuracy on HEST-1k. Several very recent methods add complementary inductive

biases: HyperST [14] aligns image and gene representations in hyperbolic space to capture their hierarchy, and Gene-DML [9] adds dual-pathway multi-level discrimination between the two modalities. We include both as feature-matched baselines. Foundation models such as STPath [15] instead pretrain across organs and sequencing platforms with explicit metadata conditioning. Most of these methods are trained and evaluated per cohort; recent benchmarks [6] emphasize the difficulty of cross-modal generalization, which motivates the cross-organ transfer protocol we additionally study. In contrast to generative prediction and large pretrained models, MorphoST uses direct regression from frozen features while retaining local and global context through distance-aware attention, at a small parameter budget.

Geometric equivariance. $E(n)$ -equivariant networks [8] use relative geometric quantities to control transformation behavior. MorphoST applies this principle to its spatial attention biases by using pairwise distances rather than absolute coordinates.

3. Method

MorphoST processes one slide at a time. The input is a set of N spots, each with a frozen UNI [1] feature $\mathbf{f}_i \in \mathbb{R}^{1024}$ and a 2D coordinate $\mathbf{p}_i \in \mathbb{R}^2$; the output is the log-normalized expression $\hat{\mathbf{y}}_i \in \mathbb{R}^G$ of a G -gene panel for every spot.

3.1. $E(2)$ -invariant spatial encoding

For each spot we build a k -nearest-neighbour graph from the Euclidean distance matrix $D_{ij} = \|\mathbf{p}_i - \mathbf{p}_j\|_2$, keeping the indices $\mathcal{N}(i)$ and distances of its k closest spots. For any orthogonal matrix R and translation $\mathbf{t}$, $\|R\mathbf{p}_i + \mathbf{t} - (R\mathbf{p}_j + \mathbf{t})\|_2 = D_{ij}$. Consequently, the spatial pathway is invariant to the 2D Euclidean group $E(2)$, which includes rotations, reflections, and translations. This statement concerns transformations of the coordinate system with the image features held fixed; it does not assume that the frozen image encoder is itself rotation invariant. We expand each distance using radial basis functions $\phi(d) = [\exp(-\gamma_m d^2)]_{m=1}^M$, where γ_m is fixed on a logarithmic grid, and use the resulting features to parameterize attention biases.

3.2. The MorphoST layer

The features are first projected, $\mathbf{x}_i = W_{\text{in}}\mathbf{f}_i$. Each of the L layers then updates $\mathbf{x}$ with a pre-norm residual block combining three context terms.

Local distance-biased attention. Each spot attends to its k neighbours with a per-head bias read out from the




Figure 2. MorphoST overview and contributions. The model receives frozen UNI morphology features and spot coordinates. Coordinates affect the network only through pairwise distances. Each layer combines local, global, and slide-level context before a linear head predicts the expression of G genes at all N spots. The design provides an invariant spatial pathway, direct single-pass inference, and a controlled framework for evaluating which forms of spatial context are useful.

radial-basis distance features:

$$\alpha_{ij} = \text{softmax}_{j \in \mathcal{N}(i)} \left(\frac{\mathbf{q}_i^\top \mathbf{k}_j}{\sqrt{d_h}} + b(\phi(D_{ij})) \right), \quad (1)$$

The distance-dependent bias allows spatial proximity to influence the attention weights while the values remain content dependent. Because the bias depends on D_{ij} rather than absolute coordinates, it preserves the coordinate invariance described above.

Global attention. A standard full self-attention over all N spots [12] captures slide-wide dependencies. This operation is permutation-equivariant with respect to the ordering of spots and does not use coordinates.

Global slide token. We summarize the slide with a content-based attention pool, $\mathbf{g} = \sum_i w_i \mathbf{x}_i$, where $w = \text{softmax}(s(\mathbf{x}))$, and broadcast a learned projection of $\mathbf{g}$ to every spot. This pathway provides an explicit summary of the slide’s overall morphological composition.

The three context terms are added to the residual stream, followed by layer normalization and a feed-forward network. A final layer normalization and linear prediction head produce expression at each spot. Pairwise-distance biases are the model’s only coordinate-dependent quantities.

3.3. Training objective and staged design

We optimize a combination of mean squared error and a correlation term,

$$\mathcal{L} = \text{MSE}(\hat{\mathbf{y}}, \mathbf{y}) + \lambda \left(1 - \frac{1}{G} \sum_g \rho_g \right), \quad (2)$$

where ρ_g denotes the Pearson correlation between predicted and measured expression for gene g across spots, and $\lambda = 0.5$. The MSE term constrains prediction scale, while the correlation term aligns optimization with the evaluation metric. To attribute performance to individual context pathways, we construct the model incrementally (Table 8): V0 an image-only MLP; V1 + global attention; V2 + local k NN attention; and V3 + the global slide token, which is the complete MorphoST configuration.

4. Experiments

4.1. Datasets

We evaluate on two public benchmarks. HEST-1k [5] provides ten tissue cohorts, each with a patient-stratified k -fold split; for presentation we aggregate the ten cohorts into eight organ groups (breast = IDC + LYMPH_IDC, colorectal = COAD + READ, others one cohort each). STImage-1K4M provides an independent multi-organ collection used to test whether conclusions transfer to a second data source. For every slide we predict the 50 most highly variable genes after $\log(1+x)$ normalization, using identical frozen UNI [1] features across all methods.

4.2. Experimental setup

Two regimes. We evaluate under (i) an intra-cohort protocol—the standard per-cohort k -fold cross-validation of STFlow [4], training separately within each cohort—and (ii) a harder cross-organ protocol

with two sample-disjoint settings: leave-one-organ-out (looo), where seven organ groups train and the eighth tests, and pooled cross-validation (pooled). Patient metadata are unavailable for some cohorts, so we do not claim the pooled splits are patient-disjoint.

Metrics and protocol. We report the Pearson correlation coefficient (PCC), Spearman correlation (SCC), and mean squared error (MSE) between predicted and measured expression, averaged over folds and three random seeds (mean_{std}). Every model selects its checkpoint on an inner validation split of the training fold and is evaluated once on the outer test fold; where a cohort fold contains a single training slide, the inner split holds out a disjoint subset of that slide’s spots. This removes the test-fold model selection present in earlier exploratory runs.

Baselines. We compare against ST-Net [3], HisToGene [7], Hist2ST [13], TRIPLEX [2], Gene-DML [9], and HyperST [14]. All methods use identical frozen UNI features, data splits, and 50-gene panels; the prediction architecture is the principal varying factor. For reference, published per-cohort accuracies of generative and higher-order methods (STFlow [4], MCToGene [11]) are summarized in Table 3; these use the ungrouped ten-cohort split and their original selection protocols and are not directly comparable to our feature-matched, nested-validation numbers.

Implementation. MorphoST uses $L=4$ layers, hidden width $d=256$, 4 attention heads, $k=8$ neighbours, and 16 radial-basis functions (4.8M parameters). We train per slide with AdamW (learning rate 10^{-3} , weight decay 10^{-2}), a cosine schedule, and gradient clipping at 1.0, for at most 100 epochs with early-stopping patience of 20; each step samples at most 3,000 spots. Experiments run on one NVIDIA A100 GPU.

4.3. Main results

Table 1 reports the intra-cohort comparison per organ, and Table 6 summarizes all methods under the three metrics. Among the feature-matched baselines evaluated under our identical nested-validation protocol, MorphoST attains the best per-organ Pearson correlation on every organ group and the best average across all three metrics (PCC 0.389, SCC 0.348, MSE 1.072; next best HyperST 0.365). Table 2 reports the same per-organ intra-organ comparison on STImage-1K4M (per-organ 5-fold, identical protocol to Table 1); runs are in progress. Table 4 reports the harder cross-organ LOOO transfer per organ on both benchmarks, and Table 5 summarizes the pooled cross-validation on both. All reported cells use inner-validation checkpoint selection rather than the earlier test-selected exploratory numbers; the corrected cross-organ baselines under this

protocol are still in progress.

Positioning against published per-cohort results. MorphoST is deliberately small (4.8M parameters) and targets cross-organ generalization rather than per-cohort state of the art. For context, Table 3 lists recently published HEST-1k accuracies, including the concurrent CVPR 2026 method MCToGene [11] and flow-based STFlow [4]. These numbers are not directly comparable to Table 1: they use the ten-cohort per-cohort split without our organ grouping (which merges the strong IDC cohort with LYMPH and COAD with READ, lowering grouped averages) and their original selection protocol rather than our nested validation. On that split larger and higher-order-interaction models score above a compact direct-regression model, as expected; our contribution is that MorphoST matches strong per-cohort baselines while being an order of magnitude smaller than TRIPLEX/HisToGene and single-step rather than iterative, and that it transfers across organs—a regime these methods do not evaluate.

4.4. Ablation: what actually matters

Table 8 reports a factorial study that isolates local attention (L), global attention (G), and the slide representation (S) rather than adding them in a fixed sequence, evaluated in the cross-organ looo regime. Global and slide context contribute most: the strongest configuration is global+slide (C011, 0.433), and every configuration containing G or S improves over the image-only baseline (C000, 0.417). Local attention is marginal for cross-organ transfer—adding it to the full model (C111, 0.421) does not improve over C011—although it is beneficial in the intra-cohort regime where the complete model is used. Table 9 separates the correlation-aware objective from the architecture and adds the independent STImage-1K4M evaluation: the correlation term raises both MorphoST and the strongest baseline by a comparable margin (+0.007 and +0.010), so it does not account for the architecture gap, and MorphoST transfers to the second benchmark (pooled PCC 0.652).

4.5. Representation and spatial-domain analysis

Beyond scalar accuracy we examine what the model represents and whether its predictions preserve tissue structure. Figure 3 visualizes learned slide representations colored by organ, probing whether the $E(2)$ -invariant pathway yields organ-transferable features. As a biology-facing check, we cluster spots by expression and compare the spatial domains obtained from predicted versus measured expression using the Adjusted Rand Index (ARI) and Normalized Mutual In-

Table 1. Intra-cohort gene expression prediction on HEST-1k, grouped by organ (Pearson $\uparrow$). Mean_{std} over three seeds with nested validation; best per row in bold. All methods use identical frozen UNI features, outer splits, 50-gene panels, and the same training budget (100 epochs, patience 20, inner-validation selection). MorphoST attains the best correlation on every organ group and on the average. Published per-cohort results, including STFlow, are in Table 3. Gene-DML [9] and HyperST [14] are our feature-matched reproductions of recent (2025–2026) methods.

Organ	ST-Net	HisToGene	Hist2ST	TRIPLEX	Gene-DML	HyperST	MorphoST (V3)
Breast	0.364 _{0.01}	0.363 _{0.04}	0.277 _{0.03}	0.399 _{0.00}	0.396 _{0.00}	0.398 _{0.00}	0.440_{0.00}
Prostate	0.251 _{0.01}	0.332 _{0.00}	0.317 _{0.02}	0.370 _{0.01}	0.363 _{0.01}	0.369 _{0.00}	0.390_{0.02}
Pancreas	0.353 _{0.02}	0.307 _{0.03}	0.243 _{0.06}	0.362 _{0.03}	0.362 _{0.04}	0.367 _{0.04}	0.392_{0.04}
Skin	0.632 _{0.01}	0.554 _{0.01}	0.509 _{0.02}	0.647 _{0.00}	0.674 _{0.00}	0.641 _{0.00}	0.696_{0.00}
Colorectal	0.181 _{0.01}	0.190 _{0.01}	0.195 _{0.01}	0.226 _{0.01}	0.222 _{0.01}	0.223 _{0.00}	0.243_{0.01}
Kidney	0.172 _{0.01}	0.210 _{0.01}	0.217 _{0.01}	0.228 _{0.03}	0.224 _{0.03}	0.256 _{0.00}	0.265_{0.02}
Liver	0.056 _{0.01}	0.029 _{0.00}	0.031 _{0.01}	0.084 _{0.00}	0.077 _{0.01}	0.087 _{0.01}	0.100_{0.00}
Lung	0.482 _{0.01}	0.490 _{0.00}	0.378 _{0.10}	0.569 _{0.00}	0.564 _{0.00}	0.581 _{0.01}	0.589_{0.01}
Average	0.311	0.309	0.271	0.361	0.360	0.365	0.389

Table 2. Intra-organ gene expression prediction on STImage-1K4M, per organ (Pearson $\uparrow$). Same protocol as Table 1: per-organ 5-fold cross-validation, Mean_{std} over three seeds with nested inner-validation selection. All methods use identical frozen UNI features, organ-specific 50-gene panels (leakage-safe: HVGs computed from training slides only), and the same training budget (100 epochs, patience 20). Corrected runs in progress.

Organ	ST-Net	HisToGene	Hist2ST	TRIPLEX	Gene-DML	HyperST	MorphoST (V3)
Breast	running	running	running	running	running	running	running
Kidney	running	running	running	running	running	running	running
Liver	running	running	running	running	running	running	running
Pancreas	running	running	running	running	running	running	running
Skin	running	running	running	running	running	running	running
Prostate	running	running	running	running	running	running	running
Brain	running	running	running	running	running	running	running
Heart	running	running	running	running	running	running	running
Average	running	running	running	running	running	running	running

Table 3. Published per-cohort accuracy on the standard splits, mean PCC $\uparrow$ as reported by MCToGene [11] (HEST-1k: ten cohorts, no organ grouping; STImage-1K4M: cancer subsets). Shown for context only—not comparable to Table 1, which uses organ grouping and nested validation. MorphoST is not evaluated on this ungrouped split.

Method	HEST-1k (10)	STImage-1K4M
ST-Net [3]	0.286	0.205
HisToGene [7]	0.237	0.232
TRIPLEX [2]	0.395	0.252
STFlow [4]	0.415	0.293
MCToGene [11]	0.435	0.316

[UMAP of learned slide representations colored by organ — generated from the completed corrected checkpoints.]

Figure 3. Learned representation. UMAP of MorphoST slide representations, colored by organ. Generated from the corrected checkpoints once training completes.

315 formation (NMI); high agreement indicates that pre-
 316 dictions recover real spatial structure rather than only
 317 per-gene trends (Table 10). All methods recover sub-
 318 stantial structure (ARI well above chance), and Mor-
 319 phoST is on par with the strongest baseline on this
 320 task, indicating that its accuracy advantage in Ta-
 321 bles 1–6 does not come at the expense of spatial co-

herence.

4.6. Qualitative results

Figure 4 overlays predicted and measured expression
 for representative marker genes on held-out slides, illus-
 trating that MorphoST reproduces spatial expression
 gradients and tissue-structured patterns rather than a

322

323

324

325

326

327

Table 4. Cross-organ LOOO transfer, per organ (Pearson $\uparrow$). Leave-one-organ-out: trained on all other organs, tested on the held-out organ. Top block: HEST-1k (8 organ groups). Bottom block: STImage-1K4M (8 organs). Mean_{std} over three seeds, nested validation. Best per row in bold. Baseline runs in progress.

Organ	ST-Net	HisToGene	Hist2ST	TRIPLEX	Gene-DML	HyperST	MorphoST (V3)
HEST-1k							
Breast	running	running	running	running	running	running	running
Prostate	running	running	running	running	running	running	running
Pancreas	running	running	running	running	running	running	running
Skin	running	running	running	running	running	running	running
Colorectal	running	running	running	running	running	running	running
Kidney	running	running	running	running	running	running	running
Liver	running	running	running	running	running	running	running
Lung	running	running	running	running	running	running	0.421 _{0.01}
Average	running	running	running	running	running	running	running
STImage-1K4M							
Breast	running	running	running	running	running	running	0.112 _{0.02}
Kidney	running	running	running	running	running	running	0.049 _{0.05}
Liver	running	running	running	running	running	running	0.178 _{0.05}
Pancreas	running	running	running	running	running	running	-0.079 _{0.05}
Skin	running	running	running	running	running	running	0.214 _{0.02}
Prostate	running	running	running	running	running	running	0.205 _{0.04}
Brain	running	running	running	running	running	running	0.043 _{0.02}
Heart	running	running	running	running	running	running	-0.034 _{0.04}
Average	running	running	running	running	running	running	0.086

Table 5. Cross-organ POOLED, average Pearson $\uparrow$ (5-fold pooled cross-validation; mean_{std} over three seeds; nested validation). Both HEST-1k and STImage-1K4M. Best per column in bold. Runs in progress.

Method	HEST-1k	STImage-1K4M
ST-Net	running	running
HisToGene	running	running
Hist2ST	running	running
TRIPLEX	running	running
Gene-DML	running	running
HyperST	running	running
MorphoST (V3)	running	0.652 _{0.01}

Table 6. Intra-cohort three-metric summary (HEST-1k, averaged over the eight organ groups; three seeds). PCC $\uparrow$, SCC $\uparrow$, MSE $\downarrow$. Best per column in bold.

Method	PCC $\uparrow$	SCC $\uparrow$	MSE $\downarrow$
ST-Net	0.311	0.269	1.633
HisToGene	0.309	0.273	1.172
Hist2ST	0.271	0.250	1.292
TRIPLEX	0.361	0.325	1.119
Gene-DML	0.360	0.323	1.113
HyperST	0.365	0.329	1.113
MorphoST (V3)	0.389	0.348	1.072

spatially uniform mean.

4.7. Efficiency and simplicity

Table 11 compares parameter counts and inference procedures. MorphoST contains 4.75M parameters, compared with 1.15M for STFlow, 13.8M for TRIPLEX, and 149M for HisToGene. Unlike STFlow’s five-step sampling procedure, MorphoST produces a slide-level prediction in one forward pass; its primary efficiency advantage is the inference procedure rather than the smallest parameter count. Table 12 reports measured model-only inference scaling.

5. Discussion

The experimental design tests two hypotheses: whether direct regression can attain accuracy comparable to flow matching under matched pathology features and splits, both within a cohort and under organ shift, and which combination of local, global, and slide-level context is most effective. Final conclusions will be drawn from the completed nested-validation runs. Independently of predictive accuracy, the architecture retains single-pass inference and coordinate-system invariance, and the representation and spatial-domain analyses probe whether its predictions are structured rather than merely well-correlated on average.

Table 7. Cross-organ generalization (three seeds, mean_{std}). looo: leave-one-organ-out transfer; pooled: pooled sample-level cross-validation. PCC $\uparrow$ / MSE $\downarrow$. Corrected runs are in progress.

Method	looo		pooled	
	PCC	MSE	PCC	MSE
ST-Net	running	running	running	running
Hist2ST	running	running	running	running
HisToGene	running	running	running	running
TRIPLEX	running	running	running	running
MorphoST (V3)	0.421 _{0.01}	1.233 _{0.11}	running	running

Table 8. Factorial context ablation (looo, Pearson mean_{std} over three seeds). L, G, and S denote local attention, global attention, and slide representation.

Variant	L	G	S	Pearson
C000				0.417 _{0.01}
C001			•	0.423 _{0.01}
C010		•		0.420 _{0.01}
C011		•	•	0.433_{0.00}
C100	•			0.396 _{0.01}
C101	•		•	0.427 _{0.00}
C110	•	•		0.416 _{0.02}
C111	•	•	•	0.421 _{0.01}

Table 9. Loss control and second benchmark. HEST-1k objective control (averaged over the eight organ groups) and STImage-1K4M replication. Pearson mean_{std} , three seeds. The correlation term helps MorphoST (+0.007) and the strongest baseline (+0.010) by a similar margin, so it does not explain the architecture gap.

Experiment	Setting	Pearson
HEST loss	MorphoST, MSE	0.382
	MorphoST, MSE+PCC	0.389
	TRIPLEX, MSE	0.361
	TRIPLEX, MSE+PCC	0.371
STImage	looo	0.086 _{0.01}
	pooled	0.652 _{0.01}

Limitations. First, MorphoST uses a correlation-aware objective, whereas the regression baselines retain MSE; Table 9 isolates this factor, but a fully controlled comparison should evaluate all architectures under both losses. Second, all methods use frozen UNI embeddings and cannot recover information omitted by the feature encoder. Third, the reported experiments use a 50-gene prediction task; larger panels and additional platforms, staining variation, and institu-

Table 10. Downstream spatial-domain recovery. Agreement between spatial domains clustered from predicted vs. measured expression on held-out slides (higher is better). ARI/NMI over held-out slides; three seeds.

Method	ARI $\uparrow$	NMI $\uparrow$
ST-Net	0.148 _{0.13}	0.207 _{0.14}
TRIPLEX	0.176 _{0.14}	0.242 _{0.15}
MorphoST (V3)	0.170 _{0.15}	0.231 _{0.15}

[Predicted vs. measured spatial expression maps for marker genes on held-out slides — generated from the saved test predictions.]

Figure 4. Predicted expression maps. Measured (top) and MorphoST-predicted (bottom) spatial expression of marker genes on held-out slides. Generated from the saved corrected test predictions.

Table 11. Model complexity. Parameter counts (baseline numbers from [4]) and inference mechanism. MorphoST predicts in a single forward pass with no generative sampling. Maximum parameter count in bold.

Model	#Params (M)	Inference
ST-Net	0.05	single pass
STFlow	1.15	5-step flow sampling
TRIPLEX	13.8	single pass
HisToGene	149.0	single pass
MorphoST (V3)	4.75	single pass

Table 12. MorphoST inference scaling on an NVIDIA A100 80GB. Median model-only latency and peak allocated GPU memory over synthetic slides; feature extraction and disk I/O are excluded.

Spots	Latency (ms)	Memory (MiB)
500	3.48	43
1,000	3.63	67
3,000	9.36	335
5,000	21.31	842
10,000	64.74	3,183

tions are needed to establish broader generalization. The coordinate-invariance result does not imply rota-

tion invariance of the frozen histology encoder. Finally, the pooled cross-organ split is sample-disjoint but cannot be verified as patient-disjoint for cohorts without patient metadata.

6. Conclusion

We introduced MorphoST, a direct-regression transformer for spatial gene expression prediction from histology. The model combines distance-biased local attention, global self-attention, and a slide-level representation; its coordinate pathway is invariant to $E(2)$ transformations. Under a unified feature-matched protocol we evaluate MorphoST in two regimes—intra-cohort and cross-organ transfer—on HEST-1k and STImage-1K4M, against direct, retrieval-based, and flow-matching baselines, with three metrics, a factorial ablation, representation and spatial-domain analyses, and measured scaling. Together these establish when direct spatial regression is a simple, reproducible, and generalizable alternative to generative prediction.

A. Numerical invariance verification

We verify the spatial-coordinate invariance independently of predictive accuracy. For a randomly initialized evaluation-mode model and a synthetic 257-spot slide, we transform only the coordinates while holding the UNI features fixed. The maximum absolute output deviations are 5.97×10^{-5} for translation, 1.97×10^{-5} for rotation, exactly zero for reflection, and 5.96×10^{-7} after permuting and restoring the spot order. These small nonzero deviations arise from floating-point distance computation. This test verifies the coordinate pathway; it does not test rotation invariance of the frozen histology encoder.

Table 13. Numerical coordinate-invariance check. Maximum absolute prediction difference.

Transformation	Max. absolute difference
Translation	5.97×10^{-5}
Rotation	1.97×10^{-5}
Reflection	0
Spot permutation	5.96×10^{-7}

B. Full per-cohort results

Table 14 reserves the full ten-cohort corrected comparison before the organ-level aggregation used in Table 1.

References

- [1] Richard J Chen, Tong Ding, Ming Y Lu, et al. Towards a general-purpose foundation model for computational pathology. *Nature Medicine*, 30:850–862, 2024. 1, 2, 3
- [2] Youngmin Chung, Ji Hun Ha, Kyeong Chan Im, and Joo Sang Lee. Accurate spatial gene expression prediction by integrating multi-resolution features. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 2, 4, 5
- [3] Bryan He, Ludvig Bergenstr hle, Linnea Stenbeck, et al. Integrating spatial gene expression and breast tumour morphology via deep learning. *Nature Biomedical Engineering*, 4(8):827–834, 2020. 2, 4, 5
- [4] Tinglin Huang, Tianyu Liu, Mehrtaash Babadi, Wengong Jin, and Rex Ying. Scalable generation of spatial transcriptomics from histology images via whole-slide flow matching. In *International Conference on Machine Learning (ICML)*, 2025. 1, 2, 3, 4, 5, 7
- [5] Guillaume Jaume, Paul Doucet, Andrew H Song, et al. HEST-1k: A dataset for spatial transcriptomics and histology image analysis. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 3
- [6] others. A large-scale benchmark of cross-modal learning for histology and gene expression in spatial transcriptomics. *arXiv preprint arXiv:2508.01490*, 2025. 2
- [7] Minoxing Pang, Kenong Su, and Mingyao Li. Leveraging information in spatial transcriptomics to predict super-resolution gene expression from histology images in tumors. *bioRxiv*, 2021. 2, 4, 5
- [8] Victor Garcia Satorras, Emiel Hoogeboom, and Max Welling. E(n) equivariant graph neural networks. In *International Conference on Machine Learning (ICML)*, 2021. 2
- [9] Yaxuan Song, Jianan Fan, Hang Chang, and Weidong Cai. Gene-dml: Dual-pathway multi-level discrimination for gene expression prediction from histopathology images. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, 2026. 2, 4, 5
- [10] Patrik L St hl, Fredrik Salm n, Sanja Vickovic, et al. Visualization and analysis of gene expression in tissue sections by spatial transcriptomics. *Science*, 353(6294):78–82, 2016. 1
- [11] Youhan Sun, Jiahua Rao, Kangrui Du, Jiancong Xie, and Yuedong Yang. Predicting spatial transcriptomics from histology images via high-order multi-cell interaction modeling. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. 2, 4, 5
- [12] Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. 3
- [13] Yuansong Zeng, Zhuoyi Wei, Weijiang Yu, et al. Spatial transcriptomics prediction from histology jointly through transformer and graph neural networks. *Briefings in Bioinformatics*, 23(5), 2022. 2, 4
- [14] Chen Zhang, Yilu An, Ying Chen, Hao Li, Xitong Ling, Lihao Liu, Junjun He, Yuxiang Lin, Zihui Wang,

Table 14. Full per-cohort results on HEST-1k (Pearson mean_{std} over three seeds; nested validation), before the organ-level aggregation of Table 1. Best per row in bold. MorphoST leads on all ten cohorts. Published per-cohort results, including STFlow, are summarized in Table 3.

Cohort	ST-Net	HisToGene	Hist2ST	TRIPLEX	Gene-DML	HyperST	MorphoST (V3)
IDC	0.538 _{0.01}	0.447 _{0.08}	0.250 _{0.07}	0.555 _{0.01}	0.557 _{0.01}	0.558 _{0.01}	0.567_{0.00}
PRAD	0.251 _{0.01}	0.332 _{0.00}	0.317 _{0.02}	0.370 _{0.01}	0.363 _{0.01}	0.369 _{0.00}	0.390_{0.02}
PAAD	0.353 _{0.02}	0.307 _{0.03}	0.243 _{0.06}	0.362 _{0.03}	0.362 _{0.04}	0.367 _{0.04}	0.392_{0.04}
SKCM	0.632 _{0.01}	0.554 _{0.01}	0.509 _{0.02}	0.647 _{0.00}	0.674 _{0.00}	0.641 _{0.00}	0.696_{0.00}
COAD	0.209 _{0.02}	0.252 _{0.01}	0.293 _{0.01}	0.293 _{0.00}	0.276 _{0.01}	0.279 _{0.01}	0.299_{0.00}
READ	0.154 _{0.00}	0.128 _{0.01}	0.097 _{0.01}	0.160 _{0.01}	0.168 _{0.01}	0.168 _{0.01}	0.187_{0.02}
CCRCC	0.172 _{0.01}	0.210 _{0.01}	0.217 _{0.01}	0.228 _{0.03}	0.224 _{0.03}	0.256 _{0.00}	0.265_{0.02}
HCC	0.056 _{0.01}	0.029 _{0.00}	0.031 _{0.01}	0.084 _{0.00}	0.077 _{0.01}	0.087 _{0.01}	0.100_{0.00}
LUNG	0.482 _{0.01}	0.490 _{0.00}	0.378 _{0.10}	0.569 _{0.00}	0.564 _{0.00}	0.581 _{0.01}	0.589_{0.01}
LYMPH_IDC	0.191 _{0.02}	0.279 _{0.01}	0.304 _{0.00}	0.242 _{0.00}	0.234 _{0.00}	0.238 _{0.01}	0.312_{0.01}

- and Rongshan Yu. Hyperst: Hierarchical hyperbolic learning for spatial transcriptomics prediction. In IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2026. 2, 4, 5
- [15] Xin Zhang et al. Stpath: a generative foundation model for integrating spatial transcriptomics and whole-slide images. npj Digital Medicine, 2025. 2