

Coordinate-Invariant Spatial Modeling for Gene Expression Prediction from Histology

Anonymous authors

Paper under double-blind review

Abstract

Predicting spatial gene expression from histology requires representations that integrate local morphology with tissue-wide context without depending on arbitrary coordinate conventions. We introduce MIST, a spatial transformer that jointly models three complementary sources of context in every layer: local attention over spatial nearest neighbors, global self-attention across all spots, and an attention-pooled slide representation broadcast to each spot. Operating on frozen histology embeddings, MIST predicts gene expression for all input spots on a slide in a single forward pass. Spatial coordinates enter only through pairwise distances, making predictions invariant to translations, rotations, and reflections of the coordinate frame when image features are fixed. We evaluate MIST on six human organ cohorts from HEST-1k under intra-cohort prediction, pooled multi-organ evaluation, and leave-one-organ-out transfer, using matched features, gene panels, splits, and model-selection procedures across methods. With combined UNI and CONCH features, MIST achieves the highest mean Pearson correlation across all six cohorts in intra-cohort evaluation, improving over the strongest baseline by 11.6% for 50-gene prediction and 11.4% for 100-gene prediction. It also improves mean correlation in pooled evaluation (0.599 versus 0.587) and transfer to held-out organs (0.422 versus 0.398). Improvements persist across UNI, CONCH, and combined features, while factorial ablations identify local spatial attention as the largest contributor to predictive accuracy. These results support explicit spatial neighborhood modeling within a unified local and tissue-wide representation, while highlighting the remaining difficulty of generalization to unseen organs.

1 Introduction

Spatial transcriptomics (ST) measures gene expression together with tissue location, enabling the study of spatial heterogeneity, tissue organization, and cellular interactions [Ståhl et al. \(2016\)](#). Its cost and experimental burden, however, constrain the number and diversity of specimens that can be profiled. In contrast, H&E slides are routinely collected and contain morphological signals associated with tissue composition and molecular state. Predicting spatial gene expression from these images could therefore extend molecular analysis to specimens without matched ST measurements.

Morphology and expression interact locally, across the slide, and at the level of slide composition. Existing predictors capture parts of this structure with patch regression, spatial graphs, retrieval, or conditional generation, but absolute coordinates can make predictions depend on an arbitrary scanner frame. We therefore seek a compact predictor that combines multi-scale context while using coordinates only through relative distances.

We introduce MIST, a direct spatial transformer that represents each spot as a token and combines distance-biased k NN attention for local context, full self-attention for slide-wide context, and attention pooling for a slide summary. Continuous radial-basis distance features preserve inter-spot geometry while guaranteeing invariance to translations, rotations, and reflections when image features are fixed (Fig. 1).

a. MIST Architecture Overview**b. Multi-context Transformer Layer**

Figure 1: MIST overview. (a) Frozen UNI, CONCH, or combined features turn H&E spot patches into tokens, while coordinate-based k NN edges provide local context. (b) Each layer combines distance-biased local attention with global self-attention before the expression head. Pairwise-distance coordinates ensure invariance to translations, rotations, and reflections.

Table 1: Architectural components across representative H&E→ST predictors. MIST is the only model that combines local neighborhood attention, slide-wide global attention, a slide-level summary token, coordinate invariance, and continuous (distance-expressive) spatial encoding rather than binary adjacency. • = present, ○ = absent.

Method	Local context	Global context	Slide token	Coord. invariant	Cont. distance
ST-Net He et al. (2020)	○	○	○	○	○
HisToGene Pang et al. (2021)	○	•	○	○	○
Hist2ST Zeng et al. (2022)	•	•	○	○	○
BLEEP Xie et al. (2023)	○	○	○	○	○
TRIPLEX Chung et al. (2024)	•	•	○	○	○
M2OST Wang et al. (2025)	•	○	○	○	○
STEM Zhu et al. (2025)	○	•	○	○	○
MIST (ours)	•	•	•	•	•

We evaluate matched features, panels, splits, budgets, and nested validation in intra-cohort, pooled cross-organ, and leave-one-organ-out (LOOO) settings. Across six human cohorts, MIST reaches mean PCC 0.496 versus 0.444 for BLEEP, 0.760 pooled PCC, and 0.490 LOOO PCC; local attention is the strongest factorial component. The model has 4.88M parameters and processes 10,000 spots in 64.7 ms on an A100.

We investigate five questions: RQ1 (Performance) accuracy across diverse organs, both per-organ and with a single pooled model; RQ2 (Transfer) generalization to an entirely unseen organ under LOOO; RQ3 (Design) the effect of encoder, context streams, and neighborhood definition; RQ4 (Structure) spatial-domain recovery; and RQ5 (Efficiency) whole-slide inference.

Contributions.

1. We introduce MIST, a multi-scale transformer with local, slide-wide, and slide-summary pathways driven by continuous pairwise-distance features and coordinate-frame invariance.
2. We quantify the effects of context, distance encoding, and feature choice through matched ablations and transformation tests.
3. We provide a controlled pooled cross-organ benchmark alongside intra-cohort and held-out-organ evaluation with matched panels, splits, baselines, and model selection.

2 Related Work

H&E→ST prediction. Prior models span patch-level regression (ST-Net [He et al. \(2020\)](#)), graph spatial context (Hist2ST [Zeng et al. \(2022\)](#)), cross-modal retrieval (BLEEP [Xie et al. \(2023\)](#)), and conditional generation (STEM [Zhu et al. \(2025\)](#)). Because they differ in features, gene panels, splits, and selection, published scores conflate data choices with architecture; we therefore compare under a matched protocol. MIST is a deterministic point predictor and does not model the full conditional expression distribution.

Geometry and generalization. Absolute position encodings can change when an equivalent tissue layout is translated, rotated, or reflected; geometric networks instead build interactions from relative quantities [Satorras et al. \(2021\)](#). MIST conditions its spatial pathway only on pairwise Euclidean distances, invariant to such frame changes for fixed histology features. For generalization we separate pooled cross-validation (organs seen, samples held out) from LOOO (an entire organ withheld) to isolate anatomical transfer on HEST-1k [Jaume et al. \(2024\)](#). An extended discussion and positioning of MIST are in Appendix B.

3 Method

MIST maps frozen spot-level histology features and measured spot coordinates to a spatial gene expression map. Its method has three stages: coordinate-invariant distance encoding, a shared multi-context layer, and supervised prediction with a correlation-aware objective.

Core Idea: MIST predicts a slide’s spatial expression in a single forward pass by fusing three views of context (local distance-biased neighborhoods, global self-attention, and a slide-level token) over frozen histology features, made invariant to the spot coordinate frame (translation, rotation, reflection) so a single model transfers across slides without per-dataset retraining.

3.1 Problem Setup

For a slide with N spots, let $\mathbf{F} = [\mathbf{f}_1, \dots, \mathbf{f}_N]^\top \in \mathbb{R}^{N \times d_{\text{in}}}$ denote frozen pathology features and $\mathbf{P} = [\mathbf{p}_1, \dots, \mathbf{p}_N]^\top \in \mathbb{R}^{N \times 2}$ the corresponding spot coordinates. The target $\mathbf{Y} \in \mathbb{R}^{N \times G}$ contains log-normalized expression for G genes. We learn $f_\theta : (\mathbf{F}, \mathbf{P}) \mapsto \hat{\mathbf{Y}}$ and process each slide independently; the model therefore supports a variable number of spots and does not require expression measurements at test time. In the feature-matched main comparisons, $\mathbf{f}_i$ concatenates frozen UNI [Chen et al. \(2024\)](#) and CONCH [Lu et al. \(2024\)](#) embeddings ($d_{\text{in}} = 1536$).

3.2 Coordinate-Invariant Distance Encoding

We first compute the pairwise Euclidean distance matrix $D_{ij} = \|\mathbf{p}_i - \mathbf{p}_j\|_2$. For each spot i , we retain the indices $\mathcal{N}(i)$ and distances of its k nearest neighbors, excluding the spot itself. Distances are expanded into M radial-basis features,

$$\phi(D_{ij}) = [\exp(-\gamma_m D_{ij}^2)]_{m=1}^M, \quad (1)$$

where the fixed bandwidth parameters γ_m are logarithmically spaced. These features influence attention weights but are never added to the spot embeddings as absolute positions.

Proposition 1 (Coordinate frame invariance):

Let $R \in \mathbb{R}^{2 \times 2}$ be orthogonal, let $\mathbf{t} \in \mathbb{R}^2$, and define $\mathbf{P}' = \mathbf{P}R^\top + \mathbf{1}\mathbf{t}^\top$. For fixed histology features $\mathbf{F}$, the prediction is unchanged:

$$f_\theta(\mathbf{F}, \mathbf{P}') = f_\theta(\mathbf{F}, \mathbf{P}). \quad (2)$$

Thus, the spatial pathway is invariant to translations, rotations, and reflections of the coordinate frame in exact arithmetic. The guarantee concerns a change of coordinate frame with fixed histology features; it does not claim that the frozen image encoder is invariant to transformations of the histology image. The proof is in Appendix I. The RBF bias is available to weight neighbors by distance, but at native HEST pixel scale the neighbor distances ($\sim 10^3$) saturate $\exp(-\gamma D^2)$ toward zero for all bandwidths, so in the reported configuration the bias contributes negligibly and the local pathway reduces to attention over the spatial k NN neighborhood. The distinguishing property of MIST is therefore this spatial locality (each spot aggregates from its geometric neighbors rather than from arbitrary or random spots), which our ablation isolates directly (spatial k NN vs. random neighbors, §5.3). The invariance guarantee above holds regardless, since it depends only on distances being rotation-invariant.

3.3 Multi-Context MIST Layer

Core Idea. At every layer, each spot receives three messages in parallel: one from its nearby spots, one from all spots, and one from a pooled summary of the slide. The model adds these messages to the current spot representation and refines the result with a feed-forward network. Only the nearby-spot message uses coordinates, and it uses pairwise distances rather than absolute positions.

The input features are projected to hidden width d , $\mathbf{X}^{(0)} = \mathbf{F}W_{\text{in}} + \mathbf{b}_{\text{in}}$. The same neighbor graph and radial-basis features are reused in all L transformer layers. At layer ℓ , a pre-normalized representation $\mathbf{H} = \text{LN}(\mathbf{X}^{(\ell)})$ is passed through three context pathways.

3.3.1 Local distance-biased attention

Each spot attends to its k neighbors with a learned, distance-dependent bias for each attention head. For head h , the attention weights are

$$\alpha_{ij}^h = \text{softmax}_{j \in \mathcal{N}(i)} \left(\frac{(\mathbf{q}_i^h)^\top \mathbf{k}_j^h}{\sqrt{d_h}} + b_h(\phi(D_{ij})) \right), \quad (3)$$

where $d_h = d/H$ for H heads. The local output is the concatenation of the weighted value sums, followed by a learned output projection. Morphological similarity is captured by the query-key term, while spatial distance contributes through the learned radial-basis bias.

3.3.2 Global self-attention

In parallel, standard multi-head self-attention Vaswani et al. (2017) allows every spot to attend to all N spots. This pathway captures long-range dependencies that may not be connected in the local graph. It uses morphology features only and is permutation-equivariant with respect to spot order.

3.3.3 Attention-pooled slide context

A learned scalar scoring function produces $a_i = \text{softmax}_i(s(\mathbf{h}_i))$. The slide representation is

$$\mathbf{g} = W_g \left(\sum_{i=1}^N a_i \mathbf{h}_i \right) + \mathbf{b}_g, \quad (4)$$

and is broadcast to every spot. Unlike a learned class token, this representation is computed directly from the slide’s spot features and summarizes its morphological composition.

For the complete model, the local (**L**), global (**G**), and broadcast slide (**S**) representations are summed and added to the residual stream. A second pre-normalization residual update applies a two-layer feed-forward network with expansion factor four and GELU activation:

$$\tilde{\mathbf{X}}^{(\ell)} = \mathbf{X}^{(\ell)} + \mathbf{L}^{(\ell)} + \mathbf{G}^{(\ell)} + \mathbf{S}^{(\ell)}, \quad (5)$$

$$\mathbf{X}^{(\ell+1)} = \tilde{\mathbf{X}}^{(\ell)} + \text{FFN}\left(\text{LN}(\tilde{\mathbf{X}}^{(\ell)})\right). \quad (6)$$

After L layers, layer normalization and a linear head map each spot representation to its G predicted gene values. The model produces the complete slide prediction in one forward pass.

Proposition 2 (Permutation equivariance):

Let $Q \in \{0, 1\}^{N \times N}$ be a permutation matrix. If the k NN construction handles distance ties equivariantly, then

$$f_{\theta}(Q\mathbf{F}, Q\mathbf{P}) = Qf_{\theta}(\mathbf{F}, \mathbf{P}). \quad (7)$$

Therefore, reordering the input spots only reorders their predicted expression profiles. Appendix I provides the proof. Our numerical checks evaluate both Propositions 1 and 2 under finite precision.

3.4 Training Objective and Ablations

For a slide with measured expression $\mathbf{Y}$, we optimize mean squared error together with a gene-wise Pearson-correlation term:

$$\mathcal{L} = \text{MSE}(\hat{\mathbf{Y}}, \mathbf{Y}) + \lambda \left(1 - \frac{1}{|\mathcal{V}|} \sum_{g \in \mathcal{V}} \rho_g \right), \quad (8)$$

where ρ_g is the Pearson correlation between predicted and measured expression across the N spots, $\mathcal{V}$ is the set of genes with nonzero target variance on that slide, and $\lambda = 0.5$. MSE constrains the scale of the predictions, while the correlation term encourages the model to reproduce spot-to-spot expression patterns; Appendix J ablates this correlation term.

The main model enables all three context pathways. Our factorial ablation independently toggles the local (L), global (G), and slide-level (S) pathways, avoiding conclusions tied to a particular order of architectural additions.

4 Experimental Setup

We evaluate on HEST-1k [Jaume et al. \(2024\)](#) with frozen UNI+CONCH patch features (Appendix D ablates each encoder in isolation) and $\log(1+x)$ targets; unless stated otherwise the target is the 50-gene HEST HVG panel. To test robustness to how the target genes are chosen, we use three 50-gene selection strategies: HVG, the highest-variance genes; DEG, differentially expressed genes selected by a Wilcoxon test across Leiden clusters; and HMHVG, a mean-and-variability criterion applied after variance filtering (Table 2). Cross-organ experiments use the six retained cohorts of Appendix A (HCC, LYMPH_IDC, and READ excluded). We study three regimes: intra-cohort (a model per organ), pooled (train across organs, hold out slide-level samples; organ identity is not a model input), and LOOO (train on all-but-one organ, test the held-out organ). Evaluation metrics. Our primary metric is the Pearson correlation coefficient (PCC) computed per gene between predicted and measured expression and averaged over genes (genes with zero target variance are excluded). We also report the mean squared error (MSE) and mean absolute error (MAE) to capture prediction scale, and, for the spatial-domain-recovery analysis, the Adjusted Rand Index (ARI) and Normalized Mutual Information (NMI) between domains clustered from predicted versus measured expression; efficiency is measured by inference latency and peak

GPU memory. All metrics are averaged over outer folds and three seeds, with checkpoints chosen on inner validation and evaluated once on the outer fold. Baselines ST-Net [He et al. \(2020\)](#), Hist2ST [Zeng et al. \(2022\)](#), BLEEP [Xie et al. \(2023\)](#), and STEM [Zhu et al. \(2025\)](#) use matched features, panels, splits, budgets, and selection. MIST has 4.88M parameters (four layers, width 256, eight neighbors). Full dataset, gene-panel, protocol, and hyperparameter details are in Appendix C.

5 Results

5.1 RQ1: Predictive performance across diverse organs

We begin by asking whether MIST predicts spatial expression accurately on organs it has seen during training. To answer this, we measure accuracy in two complementary ways under matched features, panels, and splits: a separate model trained per cohort (intra-cohort), and a single model trained jointly across all organs (pooled).

Starting with the intra-cohort setting, we stress-test accuracy along two axes of the target panel and report both. First, across the three 50-gene selection strategies (Table 2), MIST attains the best mean PCC under every strategy: 0.496 (HVG), 0.615 (DEG), and 0.650 (HMHVG), each ahead of the strongest baseline (BLEEP at 0.444, 0.559, and 0.587, respectively). Second, across panel sizes (Table 3), this lead persists at every size, as MIST tops the average at PCC@10 (0.600), PCC@50 (0.496), and PCC@100 (0.559). Focusing on the default PCC@50 panel, the gain over BLEEP is 11.6%, and MIST wins all six cohorts, with the margin widest on kidney and prostate (+0.072, +0.064) and narrowest on lung (+0.024); at the smallest 10-gene panel, correlations rise across all methods because the retained genes are the most variable, and hence the most predictable. Crucially, these gains are not an artifact of the correlation metric: as Table 4 reports, MIST also achieves the lowest error under every strategy and size (for example, at PCC@50 its MSE is 1.233 and MAE 0.794, versus 1.318 and 0.819 for BLEEP).

Table 4: Error metrics across gene-selection strategies and panel sizes on HEST-1k. Mean over three seeds, averaged over the same six cohorts as Tables 2 and 3 (IDC, CCRCC, LUNG, PAAD, PRAD, SKCM); lower values are better and the best result in each row is in bold. Models, features, splits, and training budget are matched.

Section	Panel	MSE ↓					MAE ↓				
		ST-Net	Hist2ST	BLEEP	STEM	MIST (Ours)	ST-Net	Hist2ST	BLEEP	STEM	MIST (Ours)
Gene-selection strategy	HVG	1.995	1.561	1.318	1.747	1.233	0.969	0.890	0.819	0.920	0.794
	DEG	3.336	1.757	1.557	1.958	1.477	1.457	1.021	0.945	1.054	0.928
	HMHVG	3.589	1.832	1.566	1.976	1.393	1.470	1.042	0.936	1.049	0.901
Gene-panel size	10	6.087	3.297	2.859	3.181	2.620	1.908	1.374	1.243	1.310	1.211
	50	1.995	1.561	1.318	1.747	1.233	0.969	0.890	0.819	0.920	0.794
	100	3.320	1.843	1.558	1.899	1.458	1.354	1.011	0.915	1.001	0.894

Turning next to the pooled setting, we test whether a single shared model can replace the per-organ ones. Trained jointly on all organs (leave-slide-out cross-validation; no per-organ specialisation and no organ labels), one MIST predicts across the diverse organs at once, adapting to each slide through its shared local, global, and slide-level context rather than a dedicated per-tissue model. Even so, it remains the best method: on the six-organ subset it attains the top mean PCC (0.599; Table 5), ahead of BLEEP (0.587) and Hist2ST (0.556), and reaches 0.760 on the full seven-organ HEST panel (BLEEP 0.758). Together, these two settings answer RQ1: MIST is the most accurate predictor on seen organs, whether specialised per organ or shared across all of them, so a separate model per tissue is not required.

Table 2: Intra-cohort PCC across 50-gene selection strategies. Mean_{std} over three seeds.

Panel	Cohort	ST-Net	Hist2ST	BLEEP	STEM	MIST (Ours)
HVG	IDC	0.551±0.01	0.305±0.11	0.519±0.00	0.434±0.00	0.581±0.01
	CCRCC	0.163±0.02	0.180±0.03	0.175±0.02	0.118±0.01	0.252±0.01
	LUNG	0.521±0.01	0.430±0.05	0.565±0.01	0.407±0.00	0.589±0.01
	PAAD	0.366±0.03	0.248±0.04	0.412±0.03	0.233±0.02	0.437±0.03
	PRAD	0.287±0.02	0.323±0.01	0.346±0.01	0.207±0.02	0.410±0.01
	SKCM	0.661±0.01	0.523±0.00	0.650±0.00	0.461±0.02	0.707±0.01
	Average	0.425	0.335	0.444	0.310	0.496
DEG	IDC	0.697±0.01	0.549±0.07	0.731±0.01	0.684±0.02	0.774±0.01
	CCRCC	0.312±0.01	0.335±0.01	0.304±0.00	0.253±0.01	0.402±0.01
	LUNG	0.566±0.01	0.518±0.01	0.621±0.01	0.440±0.01	0.662±0.00
	PAAD	0.425±0.03	0.352±0.01	0.487±0.03	0.306±0.04	0.518±0.03
	PRAD	0.413±0.00	0.437±0.01	0.466±0.00	0.355±0.03	0.539±0.01
	SKCM	0.752±0.01	0.656±0.00	0.746±0.03	0.617±0.02	0.794±0.00
	Average	0.528	0.475	0.559	0.443	0.615
HMHVG	IDC	0.724±0.01	0.633±0.08	0.739±0.02	0.687±0.01	0.774±0.01
	CCRCC	0.339±0.01	0.387±0.01	0.349±0.01	0.290±0.01	0.454±0.02
	LUNG	0.568±0.00	0.598±0.00	0.661±0.00	0.517±0.01	0.693±0.00
	PAAD	0.441±0.03	0.341±0.04	0.493±0.02	0.305±0.03	0.524±0.02
	PRAD	0.452±0.00	0.477±0.01	0.432±0.04	0.337±0.02	0.555±0.02
	SKCM	0.824±0.03	0.710±0.14	0.849±0.02	0.731±0.06	0.901±0.00
	Average	0.558	0.524	0.587	0.478	0.650

Table 3: Intra-cohort PCC across gene-panel sizes. Mean_{std} over three seeds; each panel reports its average.

Metric	Cohort	ST-Net	Hist2ST	BLEEP	STEM	MIST (Ours)
PCC@10	IDC	0.666±0.01	0.636±0.01	0.635±0.02	0.591±0.02	0.708±0.00
	CCRCC	0.231±0.01	0.273±0.02	0.231±0.02	0.170±0.00	0.322±0.01
	LUNG	0.702±0.01	0.519±0.02	0.735±0.00	0.616±0.01	0.759±0.00
	PAAD	0.511±0.04	0.377±0.07	0.536±0.03	0.393±0.04	0.519±0.05
	PRAD	0.442±0.03	0.459±0.02	0.434±0.03	0.325±0.03	0.524±0.02
	SKCM	0.728±0.00	0.471±0.04	0.697±0.01	0.598±0.01	0.769±0.01
	Average	0.547	0.456	0.545	0.449	0.600
PCC@50	IDC	0.551±0.01	0.309±0.07	0.519±0.00	0.439±0.01	0.581±0.01
	CCRCC	0.163±0.02	0.184±0.02	0.175±0.02	0.123±0.00	0.252±0.01
	LUNG	0.518±0.01	0.470±0.01	0.562±0.01	0.423±0.03	0.589±0.01
	PAAD	0.366±0.03	0.248±0.04	0.412±0.03	0.233±0.02	0.437±0.03
	PRAD	0.294±0.02	0.333±0.02	0.348±0.01	0.207±0.02	0.410±0.01
	SKCM	0.665±0.01	0.522±0.03	0.645±0.01	0.491±0.04	0.707±0.01
	Average	0.425	0.335	0.444	0.310	0.496
PCC@100	IDC	0.611±0.01	0.473±0.07	0.632±0.01	0.565±0.01	0.682±0.01
	CCRCC	0.221±0.01	0.274±0.03	0.247±0.02	0.192±0.00 [†]	0.328±0.02
	LUNG	0.595±0.00	0.434±0.03	0.659±0.01	0.498±0.03	0.688±0.00
	PAAD	0.355±0.04	0.222±0.04	0.399±0.04	0.236±0.02	0.435±0.03
	PRAD	0.327±0.01	0.341±0.03	0.377±0.02	0.240±0.03	0.450±0.00
	SKCM	0.696±0.01	0.446±0.12	0.696±0.01	0.602±0.02	0.771±0.01
	Average	0.468	0.365	0.502	0.389	0.559

Table 5: Cross-organ POOLED on the six-organ subset (kidney, breast, lung, pancreas, prostate, skin; COAD excluded). 5-fold pooled cross-validation (slide-level holdout, all organs seen); mean_{std} over three seeds; best in bold.

Method	PCC ↑
ST-Net He et al. (2020) (NBE’20)	0.351 _{0.02}
Hist2ST Zeng et al. (2022) (BIB’22)	0.556 _{0.01}
BLEEP Xie et al. (2023) (NeurIPS’23)	0.587 _{0.00}
STEM Zhu et al. (2025) (ICLR’25)	0.273 _{0.00}
MIST	0.599_{0.01}

Key Finding (RQ1): MIST gives the best accuracy across diverse organs both per organ (intra-cohort PCC@50 0.496 vs. 0.444 for BLEEP, winning all six cohorts) and as a single pooled model (0.599 six-organ / 0.760 full HEST), so a separate model per tissue is unnecessary when the organ is represented in training.

5.2 RQ2: Cross-organ transfer to an unseen organ

Having established strong performance on seen organs, we now ask the harder question of whether MIST transfers to an organ it has never seen. To test this, we use leave-one-organ-out (LOOO) evaluation: we train on all-but-one organ and test on the held-out one, so the target anatomy is entirely absent from training. As expected, accuracy falls relative to the pooled setting of RQ1, yet MIST still transfers best: on the six-organ subset it wins every held-out organ and attains the highest mean PCC (0.422). Figure 2 illustrates this for four representative held-out organs (kidney and prostate are omitted for readability), where MIST (orange) leads in every panel; the complete per-cohort results for all six organs are reported in the supplementary (Appendix E, Table 11). On the full seven-organ HEST panel MIST reaches 0.490 versus 0.473 for BLEEP, the strongest baseline (although BLEEP is stronger on IDC and SKCM there). Together, these results show that MIST’s inductive biases carry over to unseen anatomy, which supports RQ2; a complementary seven-organ mouse experiment in Appendix G shows the same trend.

Key Finding (RQ2): The pooled→LOOO drop, from 0.760 with the organ seen (RQ1) to 0.490 with the organ held out, isolates unseen-organ transfer as the dominant remaining gap; predicting seen organs well does not by itself transfer to a new organ.




Figure 2: Per-cohort cross-organ LOOO accuracy (four representative held-out cohorts; kidney and prostate omitted for readability, frozen UNI+CONCH). One panel per held-out cohort; box = per-gene PCC over three seeds, $\diamond$ = mean. MIST (orange) leads in all four panels, and wins all six held-out organs in the complete table (Appendix E, Table 11).

5.3 RQ3: What drives accuracy and coordinate invariance?

Given MIST’s accuracy, we next ask which components drive it and whether its coordinate invariance holds in practice. To isolate each context stream, we run a factorial study over all eight combinations of local (L), global (G), and slide-level (S) context on SKCM and LUNG. The picture is clear: local attention provides the largest gain, since enabling it alone raises mean PCC from 0.608 to 0.640, whereas global and slide context add little without it; the complete model is nonetheless best (0.643). This identifies distance-biased neighborhood aggregation as the primary contributor in the intra-cohort setting, which supports the design rationale in RQ3 (claims under organ shift would require a matched LOOO ablation, not reported here).

Beyond accuracy, we also verify the model’s coordinate invariance empirically. As Table 7 shows, prediction differences under translation and rotation remain below 6×10^{-5} , at-

Table 7: Numerical coordinate-invariance check. Maximum absolute prediction difference when only the spot coordinates are transformed (randomly initialized model, 257-spot synthetic slide).

Transformation	Max. abs. difference
Translation	5.97×10^{-5}
Rotation	1.97×10^{-5}
Reflection	0
Spot permutation	5.96×10^{-7}

tributable to floating-point evaluation, and reflection is exact. This numerically confirms the analytic guarantee, completing the answer to RQ3.

Key Finding (RQ3): Local distance-biased attention is the primary driver of accuracy (enabling it alone lifts mean PCC from 0.608 to 0.640, while global and slide context add little without it), and the predictor is provably invariant to spot translation, rotation, and reflection.

Table 6: Context-stream ablation (six-organ POOLED, no COAD). Each row toggles a subset of MIST’s local (L), global (G), and slide (S) context streams on the six-organ POOLED set (CCRCC/IDC/LUNG/PAAD/PRAD/SKCM; frozen UNI+CONCH, 14-gene, 5-fold). POOLED PCC $\uparrow$, MSE $\downarrow$, and MAE $\downarrow$ are mean_{std} over three seeds; best per column in bold. The last row replaces spatial-kNN local attention with random neighbours. The complete model (L+G+S) is best on PCC and MSE; local attention is the dominant stream.

Configuration	Streams	PCC $\uparrow$	MSE $\downarrow$	MAE $\downarrow$
Global + Local	L + G	0.60 $\pm$ 0.01	0.46 $\pm$ 0.02	0.41 $\pm$ 0.01
Global + Slide	G + S	0.57 $\pm$ 0.04	0.52 $\pm$ 0.13	0.44 $\pm$ 0.05
Global only	G	0.59 $\pm$ 0.00	0.47 $\pm$ 0.02	0.42 $\pm$ 0.01
Local + Slide	L + S	0.60 $\pm$ 0.01	0.43 $\pm$ 0.03	0.40$\pm$0.02
Local only	L	0.60 $\pm$ 0.02	0.46 $\pm$ 0.05	0.41 $\pm$ 0.02
Slide only	S	0.58 $\pm$ 0.02	0.47 $\pm$ 0.03	0.42 $\pm$ 0.02
Global + Local (random nbr.)	L _{rand} + G	0.58 $\pm$ 0.01	0.45 $\pm$ 0.01	0.40$\pm$0.00
Full, random nbr.	L _{rand} + G + S	0.58 $\pm$ 0.01	0.47 $\pm$ 0.01	0.41 $\pm$ 0.00
MIST (ours)	L + G + S	0.61$\pm$0.01	0.42$\pm$0.01	0.40$\pm$0.01



Figure 3: Qualitative spatial expression maps for a held-out skin (SKCM) slide (TENX117), three genes (rows). Each panel colors every spot by expression; MIST (right) most faithfully recovers the ground-truth spatial structure: the perivascular ring for LYZ/CD52 and the epidermal edge for KRTDAP. It attains the highest per-gene PCC (in orange), whereas STEM is noisy and Hist2ST over-smooths. Predictions use identical frozen UNI+CONCH features.

5.4 RQ4: Spatial structure preservation

Accurate per-gene prediction does not by itself guarantee that predictions preserve coherent tissue regions, so we finally ask whether MIST recovers spatial structure. To test this, we cluster spots (KMeans, $k=6$ domains) from either predicted or measured expression and compare the induced partitions using the adjusted Rand index (ARI) and normalized mutual information (NMI) on held-out slides of the six-organ subset (Table 8). Here MIST again leads, on both ARI (0.173) and NMI (0.242), ahead of the strongest baseline on each metric (BLEEP at 0.147 ARI, ST-Net at 0.201 NMI). That said, the differences are small and slide-level variation is large ($\text{std} \approx \text{mean}$), so we caution against reading them as uniformly improved domain recovery; the full seven-cohort version appears in Appendix H.

Table 8: Spatial-domain recovery (six-organ subset). ARI/NMI between domains clustered (KMeans, $k=6$) from predicted vs. measured expression on held-out slides of CCRCC, IDC, LUNG, PAAD, PRAD, and SKCM (COAD excluded); mean_{std} over three seeds, higher is better, best per column in bold.

Method	ARI $\uparrow$	NMI $\uparrow$
ST-Net He et al. (2020)	0.140 _{0.11}	0.201 _{0.12}
Hist2ST Zeng et al. (2022)	0.133 _{0.11}	0.180 _{0.11}
BLEEP Xie et al. (2023)	0.147 _{0.12}	0.197 _{0.13}
STEM Zhu et al. (2025)	0.101 _{0.12}	0.136 _{0.14}
MIST	0.173_{0.14}	0.242_{0.14}

5.5 RQ5: Computational efficiency

Finally, we ask whether this accuracy comes at a practical cost. MIST contains 4.88M parameters and predicts a complete slide in a single forward pass (Appendix Table 12). Its advantage is not the smallest footprint: ST-Net and BLEEP have fewer parameters, but MIST instead pairs multi-scale spatial context with direct, non-iterative inference. Concretely, model-only latency grows from 3.5ms for 500 spots to 64.7ms for 10,000 spots on

an A100, with 3.1 GiB peak allocated memory at the largest size (Appendix Table 13; feature extraction and disk I/O excluded). In contrast, STEM’s iterative diffusion is the only approach with substantially higher latency. MIST therefore answers RQ5 by delivering its accuracy at single-pass inference cost.

Three findings merit emphasis. First, a compact direct predictor is competitive when features, gene panels, splits, and model selection are controlled; the 10.3% relative gain over the strongest baseline rests on matched experimental conditions, not a larger model. Second, pooled multi-organ performance is not a substitute for unseen-organ evaluation: the drop from 0.760 (pooled) to 0.490 (LOOO) shows that anatomical shift remains a separate and harder problem. Third, local distance-biased attention is the dominant contributor in the factorial study; global and slide pathways provide moderate complementary gains, and the complete model (C111) achieves the best mean across both tested cohorts.

Direct and generative models serve different purposes. STEM Zhu et al. (2025) models a conditional expression distribution and can represent one-to-many uncertainty through iterative sampling. MIST produces a point estimate optimized for correlation and squared error. Its benefits, including deterministic, single-pass inference and explicit coordinate invariance, do not imply that morphology uniquely determines expression.

Limitations. The evaluation predicts 50 genes from frozen UNI+CONCH embeddings; results may not extend to transcriptome-scale output or end-to-end image encoding. LOOO uses a fixed benchmark panel and does not test platform or institution shift simultaneously. Coordinate invariance holds only for the spatial pathway with fixed image features; rotating the histology image can change its frozen embedding. Domain-recovery analysis covers only two baselines and has high slide-level variance. Pooled folds are sample-disjoint but cannot be verified as patient-disjoint where patient identifiers are unavailable.

6 Conclusion and Future Work

We developed MIST, a coordinate-invariant local-global spatial transformer for predicting gene expression from histology. Pairwise-distance encoding makes the spatial pathway invariant to rigid coordinate transformations by construction, while three parallel context signals: local k NN attention, global self-attention, and slide-level pooling, together produce complete-slide predictions in a single pass. These design choices directly address our design goals that multi-scale context is necessary and that architectural positional bias should not break under coordinate shift. Under matched HEST-1k evaluation, MIST improves mean intra-cohort correlation and remains competitive in LOOO transfer; the pooled-to-LOOO gap exposes the remaining difficulty of anatomical transfer as the central open challenge.

For future work, it would be valuable to explore how conditioning on slide-level metadata (organ type, sequencing platform) can narrow this anatomical gap, for instance, through FiLM or adaptive layer normalisation conditioners that are available at inference time without requiring expression labels. Scaling the gene panel beyond 50 genes, incorporating end-to-end image encoding, and combining MIST’s geometric structure with calibrated generative uncertainty modeling (as in STEM) are further natural directions.

AI use statement

Generative AI tools assisted with manuscript organization, language editing, ICLR formatting, and consistency checks between narrative claims and reported tables. They were not used to generate experimental data or execute model experiments. The authors reviewed all AI-assisted revisions and take responsibility for the final content of this work.

References

Richard J Chen, Tong Ding, Ming Y Lu, et al. Towards a general-purpose foundation model for computational pathology. *Nature Medicine*, 30:850–862, 2024. 3

- Youngmin Chung, Ji Hun Ha, Kyeong Chan Im, and Joo Sang Lee. Accurate spatial gene expression prediction by integrating multi-resolution features. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp. 11591–11600, 2024. doi: 10.1109/CVPR52733.2024.01101. 2
- Bryan He, Ludvig Bergenstr hle, Linnea Stenbeck, et al. Integrating spatial gene expression and breast tumour morphology via deep learning. *Nature Biomedical Engineering*, 4(8): 827–834, 2020. 2, 3, 6, 7, 10, 14, 15, 17, 18
- Guillaume Jaume, Paul Doucet, Andrew H Song, et al. HEST-1k: A dataset for spatial transcriptomics and histology image analysis. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 3, 5, 15
- Ming Y. Lu, Bowen Chen, Drew F. K. Williamson, et al. A visual-language foundation model for computational pathology. *Nature Medicine*, 30:863–874, 2024. doi: 10.1038/s41591-024-02856-4. 3
- Minxing Pang, Kenong Su, and Mingyao Li. Leveraging information in spatial transcriptomics to predict super-resolution gene expression from histology images in tumors. *bioRxiv*, 2021. doi: 10.1101/2021.11.28.470212. 2
- Victor Garcia Satorras, Emiel Hooeboom, and Max Welling. E(n) equivariant graph neural networks. In *International Conference on Machine Learning (ICML)*, 2021. 3, 14
- Patrik L St hl, Fredrik Salm n, Sanja Vickovic, et al. Visualization and analysis of gene expression in tissue sections by spatial transcriptomics. *Science*, 353(6294):78–82, 2016. 1
- Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. 4
- Hongyi Wang, Xiujun Du, Jing Liu, Shuyi Ouyang, Yen-Wei Chen, and Lanfen Lin. M2OST: Many-to-one regression for predicting spatial transcriptomics from digital pathology images. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 7709–7717, 2025. doi: 10.1609/aaai.v39i7.32830. 2
- Ronald Xie, Kuan Pang, Sai Chung, et al. Spatially resolved gene expression prediction from histology images via bi-modal contrastive learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. 2, 3, 6, 7, 10, 14, 15, 17, 18
- Yuansong Zeng, Zhuoyi Wei, Weijiang Yu, et al. Spatial transcriptomics prediction from histology jointly through transformer and graph neural networks. *Briefings in Bioinformatics*, 23(5), 2022. 2, 3, 6, 7, 10, 14, 15, 17, 18
- Sichen Zhu, Yuchen Zhu, Molei Tao, and Peng Qiu. Diffusion generative modeling for spatially resolved gene expression inference from histology images. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=FtjLUHyZAO>. 2, 3, 6, 7, 10, 11, 14, 15, 17, 18

Supplementary Material

Table of Contents

A	Dataset	14
B	Extended Related Work	14
C	Extended Experimental Setup	15
D	Encoder Ablation	15
E	Complete LOOO (Leave-One-Organ-Out) Results	16
F	Computation Resources, Model Complexity, and Inference Scaling	16
G	Mouse cross-organ diagnostic	17
H	Spatial-domain recovery	18
I	Proofs of the formal guarantees	18
I.1	Proof of Proposition 1	18
I.2	Proof of Proposition 2	18
J	Training Objective Ablation	18

A Dataset

Table 9 lists the number of samples, total spots, and cross-validation folds for each HEST-1k cohort. Two organ groups in HEST-1k contain multiple cohorts: breast (IDC + LYMPH_IDC) and colorectal (COAD + READ). Within each group the constituent cohorts differ substantially in size: IDC has $2.25\times$ more spots than LYMPH_IDC (44,974 vs. 19,964), and COAD has $2.2\times$ more spots than READ (18,523 vs. 8,407). To avoid misrepresentation from unweighted averaging over unequal cohorts, Table 3 uses IDC as the representative breast cohort and COAD as the representative colorectal cohort, discarding the smaller partner cohorts (LYMPH_IDC and READ) from the organ-level rows. This choice maximises statistical reliability of each organ-level estimate. Our cross-organ experiments (leave-one-organ-out and pooled) use the six cohorts marked in Table 9: CCRCC (kidney, 24 samples / 74,220 spots), IDC (breast, 4 / 44,974), LUNG (lung, 2 / 7,505), PAAD (pancreas, 3 / 9,793), PRAD (prostate, 23 / 62,710), and SKCM (skin, 2 / 5,716); COAD, READ, LYMPH_IDC, and HCC are not used in the cross-organ setting.

Table 9: HEST-1k cohort statistics. Spots counted after barcode filtering. Folds = number of train/test splits in the per-cohort cross-validation. The six cohorts marked $\checkmark$ (CCRCC, IDC, LUNG, PAAD, PRAD, SKCM) form the no-COAD cross-organ evaluation set used in our leave-one-organ-out and pooled experiments; LYMPH_IDC, COAD, READ, and HCC (liver) are not used there (HCC has a near-zero signal floor).

Cohort	Organ	Samples	Spots	Folds	6-organ set
IDC	Breast	4	44,974	4	$\checkmark$
LYMPH_IDC	Breast	4	19,964	4	
COAD	Colorectal	4	18,523	2	
READ	Colorectal	4	8,407	2	
CCRCC	Kidney	24	74,220	6	$\checkmark$
HCC	Liver*	2	4,248	2	
LUNG	Lung	2	7,505	2	$\checkmark$
PAAD	Pancreas	3	9,793	3	$\checkmark$
PRAD	Prostate	23	62,710	2	$\checkmark$
SKCM	Skin	2	5,716	2	$\checkmark$

*HCC (liver) is included here for completeness but excluded from Table 3; all baselines score near zero on this cohort due to low spatial gene-expression variance.

B Extended Related Work

Machine-learning prediction of gene expression from histology. ST-Net He et al. (2020) established patch-level regression. Hist2ST Zeng et al. (2022) adds spatial context through a graph over neighboring spots, while BLEEP Xie et al. (2023) aligns image and expression embeddings for cross-modal retrieval. STEM Zhu et al. (2025) instead treats expression prediction as conditional generation. These methods use different features, gene panels, splits, and selection procedures, so their published scores do not isolate architectural gains; our comparisons therefore use a matched protocol. MIST is a deterministic point predictor and does not model the full conditional expression distribution.

Geometric structure and coordinate invariance. Spatial prediction depends on the coordinate representation supplied to a model. Absolute position embeddings can encode location within a particular coordinate frame, but their output may change when an equivalent tissue layout is translated, rotated, or reflected. Geometric networks address this by constructing interactions from relative quantities Satorras et al. (2021). MIST uses pairwise Euclidean distances, which are preserved by translations, rotations, and reflections; the guarantee applies to the spatial pathway for fixed histology features, not to rotations of the image-encoder input.

Generalization across samples and organs. HEST-1k [Jaume et al. \(2024\)](#) spans many cohorts and organs. Within-cohort evaluation can exploit organ-specific morphology or acquisition effects. We therefore distinguish pooled cross-validation, which holds out samples while retaining their organ types in training, from LOOO, which withholds an entire organ and provides a stricter test of anatomical transfer.

Positioning of MIST. MIST connects these three directions: an inductive-bias-explicit spatial transformer for direct regression, evaluated under a matched cross-cohort protocol. Unlike absolute position encodings or frame-averaging strategies, its coordinate pathway is invariant by construction. Unlike generative models, it does not represent expression uncertainty; its advantage is deterministic, single-pass inference with explicit geometric structure. Unlike prior evaluations that report within-cohort accuracy only, our controlled comparison separates interpolation from anatomical transfer, making the difficulty of organ shift directly measurable.

C Extended Experimental Setup

Dataset and preprocessing. We evaluate on HEST-1k [Jaume et al. \(2024\)](#) using the cohorts of Table 9; HCC, LYMPH_IDC, and READ are excluded under the documented quality and split criteria, and the six cohorts marked in that table (CCRCC, IDC, LUNG, PAAD, PRAD, SKCM) form the cross-organ set. We use the provided patient-stratified folds, $\log(1+x)$ normalization, and frozen UNI+CONCH features (1536-d). Unless stated otherwise, targets are the 50-gene HEST HVG panel. The panel study additionally uses matched 50-gene DEG (Wilcoxon selection across Leiden clusters) and HMMHVG (mean-and-variability selection after variance filtering) panels (Table 2).

Evaluation protocol. Intra-cohort training fits a model within each cohort. Pooled five-fold validation trains across organs while holding out slide-level samples; organ identity defines the protocol but is not a model input. LOOO trains on all-but-one organ and tests the held-out organ. We report PCC, SCC, and MSE averaged over outer folds and three seeds. Checkpoints are selected on inner validation and evaluated once on the outer test fold; single-slide folds use disjoint spots for inner validation. Patient-level disjointness is not guaranteed where identifiers are unavailable.

Baselines and implementation. We compare with ST-Net [He et al. \(2020\)](#), Hist2ST [Zeng et al. \(2022\)](#), BLEEP [Xie et al. \(2023\)](#), and STEM [Zhu et al. \(2025\)](#) under matched features, panels, splits, budgets, and checkpoint selection. MIST uses four layers, width 256, four heads, eight neighbors, 16 radial bases, and 4.88M parameters. AdamW training uses learning rate 10^{-3} , weight decay 10^{-2} , cosine decay, gradient clipping, 100 epochs maximum, patience 20, and at most 3,000 spots per step.

D Encoder Ablation

We repeat the intra-cohort PCC@50 comparison (the 50-gene block of Table 3, which uses the concatenated UNI + CONCH features) with each foundation encoder in isolation: UNI only (ViT-L/16, 1024-d) and CONCH only (512-d), holding models, cohorts, gene panels, splits, and training budget fixed (Table 10). — entries are still training.

Table 10: Encoder ablation: intra-cohort PCC@50 with each foundation encoder in isolation. UNI only (1024-d) vs. CONCH only (512-d), per cohort; mean_{std} over three seeds; best per row in bold. Compare with the UNI + CONCH (1536-d) results in Table 3.

Encoder	Cohort	ST-Net	Hist2ST	BLEEP	STEM	MIST (Ours)
UNI (1024-d)	IDC Breast	0.538±0.008	0.250±0.067	0.508±0.005	—	0.569±0.009
	CCRCC Kidney	0.172±0.010	0.214±0.015	0.184±0.007	—	0.264±0.010
	LUNG Lung	0.482±0.012	0.378±0.097	0.559±0.007	0.393±0.009	0.590±0.005
	PAAD Pancreas	0.353±0.022	0.243±0.061	0.393±0.035	—	0.392±0.040
	PRAD Prostate	0.251±0.008	0.322±0.020	0.343±0.009	—	0.394±0.003
	SKCM Skin	0.632±0.010	0.510±0.015	0.639±0.015	0.471±0.031	0.694±0.001
	Average	0.405	0.319	0.438	—	0.484
CONCH (512-d)	IDC Breast	0.515±0.007	0.264±0.018	0.496±0.011	0.396±0.011	0.551±0.004
	CCRCC Kidney	0.160±0.010	0.126±0.026	0.141±0.012	—	0.222±0.010
	LUNG Lung	0.517±0.001	0.435±0.006	0.576±0.003	0.386±0.030	0.554±0.005
	PAAD Pancreas	0.372±0.035	0.186±0.035	0.384±0.033	0.251±0.028	0.387±0.045
	PRAD Prostate	0.327±0.003	0.322±0.009	0.364±0.010	0.210±0.013	0.406±0.006
	SKCM Skin	0.633±0.002	0.388±0.022	0.633±0.005	0.453±0.014	0.658±0.007
	Average	0.421	0.287	0.432	—	0.463

E Complete LOOO (Leave-One-Organ-Out) Results

Table 11 reports the complete per-cohort leave-one-organ-out transfer results on the six-organ subset summarised in the main text (Section 5.2, Figure 2). Each fold holds out one entire organ and trains on the other five, using leakage-safe per-fold gene panels and identical frozen UNI + CONCH features. MIST wins every held-out cohort and the average (0.422).

Table 11: Cross-organ LOOO transfer on a six-organ subset (kidney, breast, lung, pancreas, prostate, skin; COAD excluded), per-cohort Pearson $\uparrow$. Self-contained leave-one-organ-out: each fold holds out one organ and trains on the other five, with leakage-safe per-fold gene panels. Identical frozen UNI + CONCH features (1536-dim); mean_{std} over three seeds; best per row in bold. MIST wins every cohort and the average.

Cohort	ST-Net NBE'20	Hist2ST BIB'22	BLEEP NeurIPS'23	STEM ICLR'25 [†]	MIST (Ours)
CCRCC Kidney	0.074±0.00	0.064±0.01	0.058±0.00	0.024±0.01	0.098±0.01
IDC Breast	0.332±0.02	0.350±0.00	0.401±0.04	0.131±0.08	0.425±0.03
LUNG Lung	0.535±0.07	0.581±0.00	0.630±0.00	0.067±0.00	0.687±0.02
PAAD Pancreas	0.382±0.00	0.214±0.17	0.497±0.03	0.052±0.00	0.499±0.03
PRAD Prostate	0.078±0.01	0.062±0.01	0.091±0.00	0.016±0.00	0.105±0.01
SKCM Skin	0.500±0.09	0.694±0.02	0.708±0.01	0.299±0.00	0.720±0.03
Average	0.317	0.328	0.398	0.098	0.422

[†]STEM incomplete (1–2 of 3 seeds); values provisional.

F Computation Resources, Model Complexity, and Inference Scaling

All experiments ran on a shared GPU cluster. Feature extraction and training used NVIDIA GPUs with compute capability $\geq$ sm_80 (A100 80GB, A100-PCIe, L40s, and H200); Volta (V100) GPUs are incompatible with our CUDA-13 / PyTorch build and were excluded.

- Encoders: frozen UNI (ViT-L/16, 1024-d) and CONCH (512-d) patch features, concatenated to 1536-d; no encoder fine-tuning.
- Model: MIST has 4.88M trainable parameters (Table 12) and predicts a whole slide in a single forward pass, with no iterative sampling.

- Training: AdamW with a cosine schedule, 100 epochs with early stopping (patience 20), nested validation, and a correlation-aware objective; three seeds per configuration.
- Inference: single-pass; median model-only latency and peak memory versus slide size are reported in Table 13, with feature extraction and disk I/O excluded.

Table 12: Model complexity. Parameter counts and inference mechanisms.

Model	#Params (M)	Inference
ST-Net He et al. (2020) (NBE’20)	0.47	single-pass
Hist2ST Zeng et al. (2022) (BIB’22)	2.18	single-pass
BLEEP Xie et al. (2023) (NeurIPS’23)	0.54	single-pass
STEM Zhu et al. (2025) (ICLR’25)	34.16	iterative
MIST (V3)	4.88	single-pass

Table 13: MIST inference scaling on an NVIDIA A100 80GB. Median model-only latency and peak allocated GPU memory over synthetic slides; feature extraction and disk I/O are excluded.

Spots	Latency (ms)	Memory (MiB)
500	3.48	43
1,000	3.63	67
3,000	9.36	335
5,000	21.31	842
10,000	64.74	3,183

G Mouse cross-organ diagnostic

We evaluate 59 mouse Visium slides spanning seven organs with mouse-specific 50-gene panels. Pooled performance reaches 0.887, whereas every method remains near the floor under LOOO; we therefore do not interpret the latter as evidence of successful cross-species organ transfer.

Table 14: Cross-organ generalization on mouse HEST (three seeds, mean_{std}). looo: leave-one-organ-out transfer; pooled: five-fold pooled cross-validation. PCC $\uparrow$ on 50 mouse HVGs.

Method	looo PCC $\uparrow$	pooled PCC
ST-Net He et al. (2020)	0.013 _{0.011}	0.357 _{0.036}
Hist2ST Zeng et al. (2022)	0.018 _{0.007}	0.829 _{0.028}
BLEEP Xie et al. (2023)	-0.006 _{0.003}	0.848 _{0.012}
STEM Zhu et al. (2025)	–	0.279 _{0.009}
MIST	0.018_{0.006}	0.887_{0.005}

$\uparrow$ All looo values are near the floor because the organ-specific gene panels do not transfer reliably.

H Spatial-domain recovery

Table 15: Spatial-domain recovery (all seven cohorts). ARI/NMI between domains clustered (KMeans, $k=6$) from predicted vs. measured expression on held-out slides of all seven retained HEST-1k cohorts; mean_{std} over three seeds, higher is better, best per column in bold. Companion to the five-organ Table 8.

Method	ARI $\uparrow$	NMI $\uparrow$
ST-Net He et al. (2020)	0.149 _{0.11}	0.213 _{0.13}
Hist2ST Zeng et al. (2022)	0.149 _{0.12}	0.197 _{0.12}
BLEEP Xie et al. (2023)	0.157 _{0.12}	0.209 _{0.14}
STEM Zhu et al. (2025)	0.113 _{0.13}	0.152 _{0.15}
MIST	0.192 _{0.15}	0.259 _{0.16}

I Proofs of the formal guarantees

I.1 Proof of Proposition 1

For every pair of spots i, j , orthogonality of R gives

$$\|R\mathbf{p}_i + \mathbf{t} - (R\mathbf{p}_j + \mathbf{t})\|_2 = \|R(\mathbf{p}_i - \mathbf{p}_j)\|_2 \quad (9)$$

$$= \|\mathbf{p}_i - \mathbf{p}_j\|_2. \quad (10)$$

Hence, transforming the coordinate frame leaves the complete distance matrix unchanged. It also leaves the k NN graph, radial basis features, and distance-dependent attention biases unchanged. The input projection is unchanged because $\mathbf{F}$ is fixed. At every layer, the local pathway therefore receives the same features, neighbors, and distance biases, while the global and slide pathways do not use coordinates. Their outputs, the residual fusion, and the feed-forward update are identical. Induction over the L layers gives identical final representations and predictions, proving Proposition 1.

I.2 Proof of Proposition 2

Let $\mathbf{F}' = Q\mathbf{F}$ and $\mathbf{P}' = Q\mathbf{P}$. The pairwise distance matrix is relabelled as $\mathbf{D}' = Q\mathbf{D}Q^\top$. Under equivariant tie handling, the k NN graph and its radial basis edge features are relabelled by the same permutation. The input projection and all pointwise linear maps, layer normalizations, activations, and feed-forward networks commute with Q .

For local attention, permuting the spots only relabels each neighborhood and its attention sum, so the local output is multiplied by Q . Global self-attention is likewise permutation equivariant because its query-key score matrix is transformed by Q on both axes and its value matrix by Q . The slide attention scores are permuted, but their weighted sum is unchanged; broadcasting that same summary to all spots is equivariant. Consequently, residual fusion preserves equivariance at every layer. Induction over the layers and application of the pointwise prediction head yield $f_\theta(Q\mathbf{F}, Q\mathbf{P}) = Qf_\theta(\mathbf{F}, \mathbf{P})$, proving Proposition 2.

J Training Objective Ablation

Table 16 isolates the contribution of the Pearson correlation term in the training objective from the architectural design, confirming that the accuracy gap over baselines is not attributable to the loss function alone.

Table 16: Objective ablation (HEST-1k). Effect of adding the Pearson correlation term to the MSE objective; averaged over seven organ groups, three seeds.

Model	Objective	PCC $\uparrow$
MIST	MSE only	0.382
MIST	MSE + PCC	0.389