

# MIST: Coordinate-Invariant Spatial Modeling for Gene Expression Prediction from Histology

Anonymous authors  
Paper under double-blind review

## Abstract

Predicting spatial gene expression from routine histology images could extend molecular analysis to tissue samples without spatial transcriptomic measurements. This task requires integrating local morphology with tissue context while avoiding sensitivity to arbitrary choices of coordinate origin and orientation. We introduce MIST, a compact spatial transformer that combines three sources of context in each layer: attention over nearby spots, attention over all spots on a slide, and a pooled slide representation. The local pathway uses continuous pairwise distances to bias attention, retaining information about spatial proximity beyond a binary neighborhood graph. Because coordinates enter the model only through these distances, predictions are invariant to translations, rotations, and reflections of the coordinate frame when image features are fixed. We evaluate MIST on seven human organ cohorts from HEST-1k through intra-cohort prediction, pooled multi-organ testing on held-out samples, and leave-one-organ-out cross-organ transfer, using matched features, gene panels, splits, and model-selection procedures within each comparison. In intra-cohort prediction of 100 genes, MIST improves mean Pearson correlation by 10.0% over the strongest baseline and leads on all seven cohorts. In pooled evaluation on held-out samples from represented organs, it improves correlation by 0.3% over the strongest baseline. In a separate 50-gene cross-organ evaluation with an entire organ held out during training, MIST improves correlation by 3.6% over BLEEP, the strongest baseline considered. Across frozen UNI, CONCH, and combined UNI+CONCH features, MIST consistently achieves the highest mean correlation among the methods in the pooled encoder comparison. Factorial ablations identify local attention as the main contributor to predictive accuracy. With 4.88M parameters and model-only inference of 64.7ms for 10,000 spots on an A100, MIST supports accurate, efficient prediction across cohorts and feature encoders, while generalization to unseen organs remains a challenge.

## 1 Introduction

Spatial transcriptomics (ST) measures gene expression together with tissue location, enabling the study of spatial heterogeneity, tissue organization, and cellular interactions [Ståhl et al. \(2016\)](#). Its cost and experimental burden, however, constrain the number and diversity of specimens that can be profiled. In contrast, H&E slides are routinely collected and contain morphological signals associated with tissue composition and molecular state. Predicting spatial gene expression from these images could therefore extend molecular analysis to specimens without matched ST measurements.

The central difficulty is that morphology and expression are related at more than one level. A spot contains local visual evidence, nearby spots define its tissue neighborhood, distant regions can share structure, and the composition of the slide can shift the expected expression profile. Existing methods capture parts of this context through independent patch regression [He et al. \(2020\)](#), spatial graphs and transformers [Zeng et al. \(2022\)](#), cross-modal retrieval [Xie et al. \(2023\)](#), or conditional generation [Zhu et al. \(2025\)](#). Yet adding spatial context

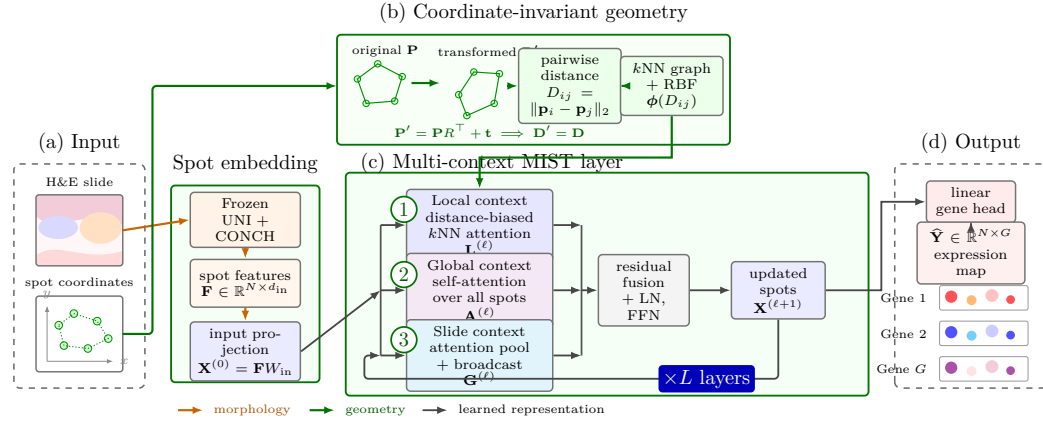


Figure 1: MIST overview. The model combines frozen spot morphology features with coordinate geometry, distance-biased local attention, global self-attention, and a pooled slide representation to predict the spatial expression map. Coordinates enter only through pairwise distances, making the spatial pathway invariant to translations, rotations, and reflections.

introduces a second concern: a model that consumes absolute positions can depend on the scanner coordinate frame. The same relative tissue arrangement may then produce different predictions after its coordinates are translated, rotated, or reflected.

This tension raises a natural question: can a compact predictor combine neighborhood context and context across the slide without becoming sensitive to an arbitrary coordinate frame? We address this question by reconsidering two common design choices.

**Coordinate representation.** Spatial models broadly fall into three groups. Methods such as Hist2ST Zeng et al. (2022) embed absolute  $x, y$  integer positions, making predictions sensitive to the scanner coordinate frame (rotation drops mean PCC from 1.00 to  $\approx 0.5$  in our experiments). Other spatial models instead build a binary  $k$ NN adjacency matrix from spot coordinates and aggregate neighbor features with uniform weights; because pairwise distances are rotation-invariant, these models are coordinate-frame invariant, but the adjacency discards the actual inter-spot distances, so all neighbors within the graph are treated as equally close regardless of how near or far they actually are. MIST takes a third path: it encodes continuous pairwise distances as radial-basis features that modulate per-head attention weights, preserving full distance information while remaining invariant to coordinate-frame changes by construction.

**Context representation.** Local aggregation, interactions across the slide, and a summary of slide composition play distinct roles. Representing them as separate pathways lets a model combine these signals while making their individual contributions directly testable through factorial ablation.

**What We Do.** We represent every measured spot as a token and let it learn from three sources in parallel: nearby spots, all spots on the slide, and a summary of the slide. Coordinates are used only to measure distances between spots. Consequently, changing the origin or orientation of the coordinate system does not change the prediction when the image features remain fixed.

Building on these observations, we introduce MIST, a compact spatial transformer for direct gene expression prediction (Fig. 1). Each layer combines three pathways:  $k$ NN attention biased by pairwise distance aggregates local neighborhoods, full self-attention models dependencies across the slide, and attention pooling constructs a slide representation that is broadcast to every spot. The three outputs are fused residually before the expression head. For fixed histology features, MIST is invariant to translations, rotations, and reflections

Table 1: Architectural components across representative H&E→ST predictors. MIST is the only model that combines local neighbourhood attention, slide-wide global attention, a slide-level summary token, coordinate invariance, and continuous (distance-expressive) spatial encoding rather than binary adjacency. • = present, ○ = absent.

| Method                                     | Local context | Global context | Slide token | Coord. invariant | Cont. distance |
|--------------------------------------------|---------------|----------------|-------------|------------------|----------------|
| ST-Net <a href="#">He et al. (2020)</a>    | ○             | ○              | ○           | ○                | ○              |
| Hist2ST <a href="#">Zeng et al. (2022)</a> | •             | •              | ○           | ○                | ○              |
| BLEEP <a href="#">Xie et al. (2023)</a>    | ○             | ○              | ○           | ○                | ○              |
| STEM <a href="#">Zhu et al. (2025)</a>     | ○             | •              | ○           | ○                | ○              |
| MIST (ours)                                | •             | •              | •           | •                | •              |

of the coordinate frame because coordinates affect the network only through pairwise distances. This guarantee concerns transformations of the coordinate representation; it does not assume that the frozen image encoder is invariant to transformations of the histology image itself.

We evaluate MIST on seven human organ cohorts from HEST-1k using matched image features, gene panels, data splits, optimization budgets, and nested validation. We distinguish prediction within individual cohorts, pooled training across represented organs, and leave one organ out (LOOO) transfer to an anatomy absent from training. Within cohorts, MIST achieves a mean Pearson correlation of 0.470, compared with 0.426 for the strongest baseline (BLEEP), and leads on all seven cohorts. It reaches 0.760 in pooled evaluation and 0.490 under LOOO transfer. A factorial ablation on SKCM and LUNG identifies local attention as the dominant contributor, while the complete model achieves the highest mean. MIST uses 4.88M parameters and predicts a slide with 10,000 spots in 64.7 ms on an A100.

We investigate five research questions: RQ1 (Accuracy): Does MIST match or exceed existing predictors for intra-cohort H&E→ST prediction across gene-panel sizes (10/50/100) and selection strategies (HVG, DEG, HMHVG)? RQ2 (Transfer): How well does MIST generalize across organs under pooled cross-validation and the stricter leave-one-organ-out (LOOO) setting? RQ3 (What matters): Which factors drive accuracy—the frozen patch encoder, the local/global/slide-level context streams, or the spatial-neighbourhood definition? RQ4 (Structure): Do MIST’s predictions preserve coherent spatial tissue structure usable for downstream analysis? RQ5 (Efficiency): Is MIST computationally efficient for whole-slide prediction?

Contributions. This work makes three contributions:

1. We introduce MIST, a multi-scale spatial transformer that fuses local, slide-wide, and slide-summary context in every layer using continuous pairwise-distance RBF features as its sole coordinate signal, providing coordinate-frame invariance while retaining distance information.
2. We establish the value of this design through rotation-sensitivity tests, factorial context ablations, and distance-encoding controls, identifying the components that drive accuracy.
3. To the best of our knowledge, we provide the first controlled pooled cross-organ evaluation for histology-to-spatial-expression prediction, alongside intra-cohort and held-out-organ tests, with matched features, gene panels, splits, baselines, and model-selection procedures.

## 2 Related Work

Machine Learning Prediction of Gene Expression from Histology. ST-Net [He et al. \(2020\)](#) established patch-level regression. Hist2ST [Zeng et al. \(2022\)](#) adds spatial context through a

graph over neighbouring spots, while BLEEP [Xie et al. \(2023\)](#) aligns image and expression embeddings for cross-modal retrieval. STEM [Zhu et al. \(2025\)](#) instead treats expression prediction as conditional generation. These methods use different features, gene panels, splits, and selection procedures, so their published scores do not isolate architectural gains; our comparisons therefore use a matched protocol. MIST is a deterministic point predictor and does not model the full conditional expression distribution.

**Geometric Structure and Coordinate Invariance.** Spatial prediction depends on the coordinate representation supplied to a model. Absolute position embeddings can encode location within a particular coordinate frame, but their output may change when an equivalent tissue layout is translated, rotated, or reflected. Geometric networks address this issue by constructing interactions from relative quantities [Satorras et al. \(2021\)](#). MIST uses pairwise Euclidean distances, which are preserved by translations, rotations, and reflections. The guarantee applies to the spatial pathway for fixed histology features, not to rotations of the image encoder input.

**Generalization across Samples and Organs.** HEST-1k [Jaume et al. \(2024\)](#) spans many cohorts and organs, but within-cohort evaluation can be organized within cohorts. A model can therefore exploit organ-specific morphology or acquisition effects. We therefore distinguish pooled cross-validation, which holds out samples while retaining their organ types in training, from LOOO, which withholds an entire organ and provides a stricter test of anatomical transfer.

**Positioning of MIST.** MIST connects these three research directions: it is an inductive-bias-explicit spatial transformer for direct regression, evaluated under a matched cross-cohort protocol. Unlike absolute position encodings or frame-averaging strategies, its coordinate pathway is invariant by construction. Unlike generative models, it does not represent expression uncertainty; its advantage is deterministic, single-pass inference with explicit geometric structure. Unlike prior evaluations that report within-cohort accuracy only, our controlled comparison separates interpolation from anatomical transfer, making the difficulty of organ shift directly measurable.

### 3 Method

MIST maps frozen spot-level histology features and measured spot coordinates to a spatial gene expression map. Its method has three stages: coordinate-invariant distance encoding, a shared multi-context layer, and supervised prediction with a correlation-aware objective.

#### 3.1 Problem Setup

For a slide with  $N$  spots, let  $\mathbf{F} = [\mathbf{f}_1, \dots, \mathbf{f}_N]^\top \in \mathbb{R}^{N \times d_{\text{in}}}$  denote frozen pathology features and  $\mathbf{P} = [\mathbf{p}_1, \dots, \mathbf{p}_N]^\top \in \mathbb{R}^{N \times 2}$  the corresponding spot coordinates. The target  $\mathbf{Y} \in \mathbb{R}^{N \times G}$  contains log-normalized expression for  $G$  genes. We learn  $f_\theta : (\mathbf{F}, \mathbf{P}) \mapsto \hat{\mathbf{Y}}$  and process each slide independently; the model therefore supports a variable number of spots and requires no expression measurements at test time. In the feature-matched main comparisons,  $\mathbf{f}_i$  concatenates frozen UNI [Chen et al. \(2024\)](#) and CONCH [Lu et al. \(2024\)](#) embeddings ( $d_{\text{in}} = 1536$ ).

#### 3.2 Coordinate-Invariant Distance Encoding

We first compute the pairwise Euclidean distance matrix  $D_{ij} = \|\mathbf{p}_i - \mathbf{p}_j\|_2$ . For each spot  $i$ , we retain the indices  $\mathcal{N}(i)$  and distances of its  $k$  nearest neighbours, excluding the spot itself. Distances are expanded into  $M$  radial-basis features,

$$\phi(D_{ij}) = [\exp(-\gamma_m D_{ij}^2)]_{m=1}^M, \quad (1)$$

where the fixed bandwidth parameters  $\gamma_m$  are logarithmically spaced. These features influence attention weights but are never added to the spot embeddings as absolute positions.

Proposition 1 (Coordinate frame invariance):

Let  $R \in \mathbb{R}^{2 \times 2}$  be orthogonal, let  $\mathbf{t} \in \mathbb{R}^2$ , and define  $\mathbf{P}' = \mathbf{P}R^\top + \mathbf{1}\mathbf{t}^\top$ . For fixed histology features  $\mathbf{F}$ , the prediction is unchanged:

$$f_\theta(\mathbf{F}, \mathbf{P}') = f_\theta(\mathbf{F}, \mathbf{P}). \quad (2)$$

Thus, the spatial pathway is invariant to translations, rotations, and reflections of the coordinate frame in exact arithmetic. The guarantee concerns a change of coordinate frame with fixed histology features; it does not claim that the frozen image encoder is invariant to transformations of the histology image. The proof is in Appendix I. Note that binary  $k$ NN adjacency models satisfy the same invariance property, since pairwise distances are also rotation-invariant. The distinguishing property of MIST is that it retains the magnitude of those distances as continuous RBF features influencing attention weights, whereas binary-adjacency models collapse each distance to a 0/1 neighbor indicator, losing the information of how close or far each neighbor actually is.

### 3.3 Multi-Context MIST Layer

**Core Idea.** At every layer, each spot receives three messages in parallel: one from its nearby spots, one from all spots, and one from a pooled summary of the slide. The model adds these messages to the current spot representation and refines the result with a feed-forward network. Only the nearby-spot message uses coordinates, and it uses pairwise distances rather than absolute positions.

The input features are projected to hidden width  $d$ ,  $\mathbf{X}^{(0)} = \mathbf{F}W_{\text{in}} + \mathbf{b}_{\text{in}}$ . The same neighbour graph and radial-basis features are reused in all  $L$  transformer layers. At layer  $\ell$ , a pre-normalized representation  $\mathbf{H} = \text{LN}(\mathbf{X}^{(\ell)})$  is passed through three context pathways.

#### 3.3.1 Local distance-biased attention

Each spot attends to its  $k$  neighbours with a learned, distance-dependent bias for each attention head. For head  $h$ , the attention weights are

$$\alpha_{ij}^h = \text{softmax}_{j \in \mathcal{N}(i)} \left( \frac{(\mathbf{q}_i^h)^\top \mathbf{k}_j^h}{\sqrt{d_h}} + b_h(\phi(D_{ij})) \right), \quad (3)$$

where  $d_h = d/H$  for  $H$  heads. The local output is the concatenation of the weighted value sums, followed by a learned output projection. Morphological similarity is captured by the query-key term, while spatial distance contributes through the learned radial-basis bias.

#### 3.3.2 Global self-attention

In parallel, standard multi-head self-attention Vaswani et al. (2017) allows every spot to attend to all  $N$  spots. This pathway captures long-range dependencies that may not be connected in the local graph. It uses morphology features only and is permutation-equivariant with respect to spot order.

#### 3.3.3 Attention-pooled slide context

A learned scalar scoring function produces  $a_i = \text{softmax}_i(s(\mathbf{h}_i))$ . The slide representation is

$$\mathbf{g} = W_g \left( \sum_{i=1}^N a_i \mathbf{h}_i \right) + \mathbf{b}_g, \quad (4)$$

and is broadcast to every spot. Unlike a learned class token, this representation is computed directly from the slide’s spot features and summarizes its morphological composition.

For the complete V3 model, the local output, global output, and broadcast slide representation are summed and added to the residual stream. A second pre-normalization residual update applies a two-layer feed-forward network with expansion factor four and GELU activation:

$$\tilde{\mathbf{X}}^{(\ell)} = \mathbf{X}^{(\ell)} + \mathbf{L}^{(\ell)} + \mathbf{A}^{(\ell)} + \mathbf{G}^{(\ell)}, \quad (5)$$

$$\mathbf{X}^{(\ell+1)} = \tilde{\mathbf{X}}^{(\ell)} + \text{FFN}\left(\text{LN}(\tilde{\mathbf{X}}^{(\ell)})\right). \quad (6)$$

After  $L$  layers, layer normalization and a linear head map each spot representation to its  $G$  predicted gene values. The model produces the complete slide prediction in one forward pass.

Proposition 2 (Permutation equivariance):

Let  $Q \in \{0, 1\}^{N \times N}$  be a permutation matrix. If the  $k$ NN construction handles distance ties equivariantly, then

$$f_{\theta}(Q\mathbf{F}, Q\mathbf{P}) = Qf_{\theta}(\mathbf{F}, \mathbf{P}). \quad (7)$$

Therefore, reordering the input spots only reorders their predicted expression profiles. Appendix I provides the proof. Our numerical checks evaluate both Propositions 1 and 2 under finite precision.

### 3.4 Training Objective and Ablations

For a slide with measured expression  $\mathbf{Y}$ , we optimize mean squared error together with a gene-wise Pearson-correlation term:

$$\mathcal{L} = \text{MSE}(\hat{\mathbf{Y}}, \mathbf{Y}) + \lambda \left( 1 - \frac{1}{|\mathcal{V}|} \sum_{g \in \mathcal{V}} \rho_g \right), \quad (8)$$

where  $\rho_g$  is the Pearson correlation between predicted and measured expression across the  $N$  spots,  $\mathcal{V}$  is the set of genes with nonzero target variance on that slide, and  $\lambda = 0.5$ . MSE constrains the scale of the predictions, while the correlation term encourages the model to reproduce spot-to-spot expression patterns.

The main model enables all three context pathways. Our factorial ablation independently toggles the local (L), global (G), and slide-level (S) pathways, avoiding conclusions tied to a particular order of architectural additions (Table 6).

## 4 Experimental Results

**Dataset and Preprocessing.** We evaluate on HEST-1k [Jaume et al. \(2024\)](#). The controlled human-tissue benchmark retains seven cohorts representing seven organs; HCC, LYMPH\_IDC, and READ are excluded for the data-quality and split constraints documented in Appendix H. Each cohort uses the provided patient-stratified folds. We predict the 50 most highly variable genes (HVGs)—genes ranked by cross-spot expression variance—after  $\log(1+x)$  normalization, using the same frozen UNI+CONCH features for every method in the main comparison. HVGs are the HEST-1k default panel and capture broad tissue-level expression variation.

**Gene selection.** Published methods use different target-gene selection protocols. ST-Net selects 250 genes with the highest mean expression (highly expressed genes, HEGs) [He et al. \(2020\)](#). Hist2ST selects the top 1,000 HVGs per section before expression-frequency filtering [Zeng et al. \(2022\)](#). BLEEP trains on a union of per-slide HVGs and reports results on HVG, HEG, and marker-gene subsets [Xie et al. \(2023\)](#). STEM evaluates high-mean highly variable gene (HMHVG) and HVG panels on HEST, and HMHVG and differentially expressed gene (DEG) panels on HER2ST; its DEGs are selected across annotated tissue regions in training slides [Zhu et al. \(2025\)](#). Thus, HMHVG in these papers denotes selection by expression mean and variability.

Table 2: Intra-cohort gene expression prediction on HEST-1k across gene-panel sizes (Pearson  $\uparrow$ ). PCC@ $k$  denotes prediction with a  $k$ -gene panel. Mean<sub>std</sub> over three seeds with nested-validation; best per cohort and panel in bold. The bottom blocks report the cohort-averaged MSE $\downarrow$  and MAE $\downarrow$  at each panel size (best, i.e. lowest, in bold). All methods use identical frozen UNI + CONCH features (1536-dim), splits, and training budget (100 epochs, patience 20, inner-validation selection). The 10-gene panel is the highest-variance subset of each cohort’s HEST 50-HVG panel; the 100-gene panel extends it with the next 50 highest-variance genes. HCC, LYMPH\_IDC, and READ are excluded (Appendix H).

| Metric               | Cohort          | ST-Net           | NBE’20                           | Hist2ST                          | BIB’22                        | BLEEP                            | NeurIPS’23 | STEM | ICLR’25 | MIST (Ours) |
|----------------------|-----------------|------------------|----------------------------------|----------------------------------|-------------------------------|----------------------------------|------------|------|---------|-------------|
| PCC@10               | IDC Breast      | 0.666 $\pm$ 0.01 | 0.636 $\pm$ 0.01                 | 0.635 $\pm$ 0.02                 | 0.591 $\pm$ 0.02              | <b>0.708<math>\pm</math>0.00</b> |            |      |         |             |
|                      | COAD Colorectal | 0.255 $\pm$ 0.02 | <b>0.447<math>\pm</math>0.02</b> | 0.413 $\pm$ 0.01                 | 0.266 $\pm$ 0.01              | 0.382 $\pm$ 0.02                 |            |      |         |             |
|                      | CCRCC Kidney    | 0.231 $\pm$ 0.01 | 0.273 $\pm$ 0.02                 | 0.231 $\pm$ 0.02                 | 0.170 $\pm$ 0.00              | <b>0.322<math>\pm</math>0.01</b> |            |      |         |             |
|                      | LUNG Lung       | 0.702 $\pm$ 0.01 | 0.519 $\pm$ 0.02                 | 0.735 $\pm$ 0.00                 | 0.616 $\pm$ 0.01              | <b>0.759<math>\pm</math>0.00</b> |            |      |         |             |
|                      | PAAD Pancreas   | 0.511 $\pm$ 0.04 | 0.377 $\pm$ 0.07                 | <b>0.536<math>\pm</math>0.03</b> | 0.393 $\pm$ 0.04              | 0.519 $\pm$ 0.05                 |            |      |         |             |
|                      | PRAD Prostate   | 0.442 $\pm$ 0.03 | 0.459 $\pm$ 0.02                 | 0.434 $\pm$ 0.03                 | 0.325 $\pm$ 0.03              | <b>0.524<math>\pm</math>0.02</b> |            |      |         |             |
|                      | SKCM Skin       | 0.728 $\pm$ 0.00 | 0.471 $\pm$ 0.04                 | 0.697 $\pm$ 0.01                 | 0.598 $\pm$ 0.01              | <b>0.769<math>\pm</math>0.01</b> |            |      |         |             |
|                      | Average         | 0.505            | 0.455                            | 0.526                            | 0.423                         | <b>0.569</b>                     |            |      |         |             |
| PCC@50               | IDC Breast      | 0.551 $\pm$ 0.01 | 0.309 $\pm$ 0.07                 | 0.519 $\pm$ 0.00                 | 0.439 $\pm$ 0.01              | <b>0.581<math>\pm</math>0.01</b> |            |      |         |             |
|                      | COAD Colorectal | 0.242 $\pm$ 0.01 | 0.295 $\pm$ 0.01                 | 0.314 $\pm$ 0.01                 | 0.177 $\pm$ 0.01              | <b>0.318<math>\pm</math>0.01</b> |            |      |         |             |
|                      | CCRCC Kidney    | 0.163 $\pm$ 0.02 | 0.184 $\pm$ 0.02                 | 0.175 $\pm$ 0.02                 | 0.123 $\pm$ 0.00              | <b>0.252<math>\pm</math>0.01</b> |            |      |         |             |
|                      | LUNG Lung       | 0.518 $\pm$ 0.01 | 0.470 $\pm$ 0.01                 | 0.562 $\pm$ 0.01                 | 0.423 $\pm$ 0.03              | <b>0.589<math>\pm</math>0.01</b> |            |      |         |             |
|                      | PAAD Pancreas   | 0.359 $\pm$ 0.03 | 0.247 $\pm$ 0.02                 | 0.419 $\pm$ 0.03                 | 0.233 $\pm$ 0.02              | <b>0.437<math>\pm</math>0.03</b> |            |      |         |             |
|                      | PRAD Prostate   | 0.294 $\pm$ 0.02 | 0.333 $\pm$ 0.02                 | 0.348 $\pm$ 0.01                 | 0.207 $\pm$ 0.02              | <b>0.410<math>\pm</math>0.01</b> |            |      |         |             |
|                      | SKCM Skin       | 0.665 $\pm$ 0.01 | 0.522 $\pm$ 0.03                 | 0.645 $\pm$ 0.01                 | 0.491 $\pm$ 0.04              | <b>0.707<math>\pm</math>0.01</b> |            |      |         |             |
|                      | Average         | 0.399            | 0.337                            | 0.426                            | 0.299                         | <b>0.470</b>                     |            |      |         |             |
| PCC@100              | IDC Breast      | 0.611 $\pm$ 0.01 | 0.473 $\pm$ 0.07                 | 0.632 $\pm$ 0.01                 | 0.565 $\pm$ 0.01              | <b>0.682<math>\pm</math>0.01</b> |            |      |         |             |
|                      | COAD Colorectal | 0.336 $\pm$ 0.00 | 0.406 $\pm$ 0.01                 | 0.488 $\pm$ 0.01                 | 0.348 $\pm$ 0.02              | <b>0.499<math>\pm</math>0.01</b> |            |      |         |             |
|                      | CCRCC Kidney    | 0.221 $\pm$ 0.01 | 0.274 $\pm$ 0.03                 | 0.247 $\pm$ 0.02                 | 0.192 $\pm$ 0.00 <sup>†</sup> | <b>0.328<math>\pm</math>0.02</b> |            |      |         |             |
|                      | LUNG Lung       | 0.595 $\pm$ 0.00 | 0.434 $\pm$ 0.03                 | 0.659 $\pm$ 0.01                 | 0.498 $\pm$ 0.03              | <b>0.688<math>\pm</math>0.00</b> |            |      |         |             |
|                      | PAAD Pancreas   | 0.355 $\pm$ 0.04 | 0.222 $\pm$ 0.04                 | 0.399 $\pm$ 0.04                 | 0.236 $\pm$ 0.02              | <b>0.435<math>\pm</math>0.03</b> |            |      |         |             |
|                      | PRAD Prostate   | 0.327 $\pm$ 0.01 | 0.341 $\pm$ 0.03                 | 0.377 $\pm$ 0.02                 | 0.240 $\pm$ 0.03              | <b>0.450<math>\pm</math>0.00</b> |            |      |         |             |
|                      | SKCM Skin       | 0.696 $\pm$ 0.01 | 0.446 $\pm$ 0.12                 | 0.696 $\pm$ 0.01                 | 0.602 $\pm$ 0.02              | <b>0.771<math>\pm</math>0.01</b> |            |      |         |             |
|                      | Average         | 0.449            | 0.371                            | 0.500                            | 0.415 <sup>†</sup>            | <b>0.550</b>                     |            |      |         |             |
| Avg MSE $\downarrow$ | @10             | 6.119            | 3.653                            | 3.334                            | 3.647                         | <b>3.105</b>                     |            |      |         |             |
|                      | @50             | 2.041            | 1.653                            | 1.445                            | 1.874                         | <b>1.388</b>                     |            |      |         |             |
|                      | @100            | 3.650            | 1.943                            | 1.690                            | 2.066                         | <b>1.606</b>                     |            |      |         |             |
| Avg MAE $\downarrow$ | @10             | 1.897            | 1.439                            | 1.332                            | 1.397                         | <b>1.306</b>                     |            |      |         |             |
|                      | @50             | 0.972            | 0.907                            | 0.843                            | 0.944                         | <b>0.829</b>                     |            |      |         |             |
|                      | @100            | 1.411            | 1.034                            | 0.948                            | 1.040                         | <b>0.929</b>                     |            |      |         |             |

STEM: corrected DiT ( $T = 1000$ , 34M params). <sup>†</sup>STEM CCRCC at PCC@100 has two of three seeds; its average excludes the pending seed.

Table 3: Intra-cohort gene expression prediction on HEST-1k across gene-selection strategies (Pearson  $\uparrow$ ). HVG: highly variable genes (the HEST default); DEG: differentially expressed genes (Leiden clustering + Wilcoxon rank-sum); HMHVG: high-mean highly variable genes selected from variance-filtered candidates using expression statistics. Fifty genes each; mean<sub>std</sub> over three seeds with nested-validation; best per cohort and strategy in bold. Same features, splits, and training budget as Table 2.

| Panel | Cohort          | ST-Net NBE'20    | Hist2ST BIB'22   | BLEEP NeurIPS'23                 | STEM ICLR'25     | MIST (Ours)                      |
|-------|-----------------|------------------|------------------|----------------------------------|------------------|----------------------------------|
| HVG   | IDC Breast      | 0.551 $\pm$ 0.01 | 0.305 $\pm$ 0.11 | 0.519 $\pm$ 0.00                 | 0.434 $\pm$ 0.00 | <b>0.581<math>\pm</math>0.01</b> |
|       | COAD Colorectal | 0.235 $\pm$ 0.01 | 0.298 $\pm$ 0.01 | <b>0.322<math>\pm</math>0.01</b> | 0.177 $\pm$ 0.01 | 0.318 $\pm$ 0.01                 |
|       | CCRCC Kidney    | 0.163 $\pm$ 0.02 | 0.180 $\pm$ 0.03 | 0.175 $\pm$ 0.02                 | 0.118 $\pm$ 0.01 | <b>0.252<math>\pm</math>0.01</b> |
|       | LUNG Lung       | 0.521 $\pm$ 0.01 | 0.430 $\pm$ 0.05 | 0.565 $\pm$ 0.01                 | 0.407 $\pm$ 0.00 | <b>0.589<math>\pm</math>0.01</b> |
|       | PAAD Pancreas   | 0.366 $\pm$ 0.03 | 0.248 $\pm$ 0.04 | 0.412 $\pm$ 0.03                 | 0.233 $\pm$ 0.02 | <b>0.437<math>\pm</math>0.03</b> |
|       | PRAD Prostate   | 0.287 $\pm$ 0.02 | 0.323 $\pm$ 0.01 | 0.346 $\pm$ 0.01                 | 0.207 $\pm$ 0.02 | <b>0.410<math>\pm</math>0.01</b> |
|       | SKCM Skin       | 0.661 $\pm$ 0.01 | 0.523 $\pm$ 0.00 | 0.650 $\pm$ 0.00                 | 0.461 $\pm$ 0.02 | <b>0.707<math>\pm</math>0.01</b> |
|       | Average         | 0.398            | 0.330            | 0.427                            | 0.291            | <b>0.471</b>                     |
| DEG   | IDC Breast      | 0.697 $\pm$ 0.01 | 0.549 $\pm$ 0.07 | 0.731 $\pm$ 0.01                 | 0.684 $\pm$ 0.02 | <b>0.774<math>\pm</math>0.01</b> |
|       | COAD Colorectal | 0.497 $\pm$ 0.00 | 0.608 $\pm$ 0.01 | 0.639 $\pm$ 0.02                 | 0.534 $\pm$ 0.01 | <b>0.677<math>\pm</math>0.00</b> |
|       | CCRCC Kidney    | 0.312 $\pm$ 0.01 | 0.335 $\pm$ 0.01 | 0.304 $\pm$ 0.00                 | 0.253 $\pm$ 0.01 | <b>0.402<math>\pm</math>0.01</b> |
|       | LUNG Lung       | 0.566 $\pm$ 0.01 | 0.518 $\pm$ 0.01 | 0.621 $\pm$ 0.01                 | 0.440 $\pm$ 0.01 | <b>0.662<math>\pm</math>0.00</b> |
|       | PAAD Pancreas   | 0.425 $\pm$ 0.03 | 0.352 $\pm$ 0.01 | 0.487 $\pm$ 0.03                 | 0.306 $\pm$ 0.04 | <b>0.518<math>\pm</math>0.03</b> |
|       | PRAD Prostate   | 0.413 $\pm$ 0.00 | 0.437 $\pm$ 0.01 | 0.466 $\pm$ 0.00                 | 0.355 $\pm$ 0.03 | <b>0.539<math>\pm</math>0.01</b> |
|       | SKCM Skin       | 0.752 $\pm$ 0.01 | 0.656 $\pm$ 0.00 | 0.746 $\pm$ 0.03                 | 0.617 $\pm$ 0.02 | <b>0.794<math>\pm</math>0.00</b> |
|       | Average         | 0.523            | 0.494            | 0.571                            | 0.456            | <b>0.624</b>                     |
| HMHVG | IDC Breast      | 0.724 $\pm$ 0.01 | 0.633 $\pm$ 0.08 | 0.739 $\pm$ 0.02                 | 0.687 $\pm$ 0.01 | <b>0.774<math>\pm</math>0.01</b> |
|       | COAD Colorectal | 0.503 $\pm$ 0.01 | 0.659 $\pm$ 0.02 | 0.645 $\pm$ 0.00                 | 0.528 $\pm$ 0.03 | <b>0.698<math>\pm</math>0.01</b> |
|       | CCRCC Kidney    | 0.339 $\pm$ 0.01 | 0.387 $\pm$ 0.01 | 0.349 $\pm$ 0.01                 | 0.290 $\pm$ 0.01 | <b>0.454<math>\pm</math>0.02</b> |
|       | LUNG Lung       | 0.568 $\pm$ 0.00 | 0.598 $\pm$ 0.00 | 0.661 $\pm$ 0.00                 | 0.517 $\pm$ 0.01 | <b>0.693<math>\pm</math>0.00</b> |
|       | PAAD Pancreas   | 0.441 $\pm$ 0.03 | 0.341 $\pm$ 0.04 | 0.493 $\pm$ 0.02                 | 0.305 $\pm$ 0.03 | <b>0.524<math>\pm</math>0.02</b> |
|       | PRAD Prostate   | 0.452 $\pm$ 0.00 | 0.477 $\pm$ 0.01 | 0.432 $\pm$ 0.04                 | 0.337 $\pm$ 0.02 | <b>0.555<math>\pm</math>0.02</b> |
|       | SKCM Skin       | 0.824 $\pm$ 0.03 | 0.710 $\pm$ 0.14 | 0.849 $\pm$ 0.02                 | 0.731 $\pm$ 0.06 | <b>0.901<math>\pm</math>0.00</b> |
|       | Average         | 0.550            | 0.544            | 0.595                            | 0.485            | <b>0.657</b>                     |

All strategies select 50 genes per cohort. MIST attains the highest average under every strategy; absolute accuracy rises HVG<DEG<HMHVG for all methods, as HMHVG panels are the most predictable from H&E.

Table 4: Cross-organ LOOO transfer, per-cohort Pearson  $\uparrow$  (HEST-1k). Mean<sub>std</sub> over three seeds; nested-validation. Leave-one-organ-out: trained on all other organ groups, tested on the held-out cohort. All methods use identical frozen UNI + CONCH features (1536-dim) and 50-gene panels. Best per row in bold.

| Cohort          | ST-Net NBE'20    | Hist2ST BIB'22   | BLEEP NeurIPS'23                 | STEM ICLR'25     | MIST+CONCH                       |
|-----------------|------------------|------------------|----------------------------------|------------------|----------------------------------|
| IDC Breast      | 0.198 $\pm$ 0.05 | 0.458 $\pm$ 0.01 | <b>0.529<math>\pm</math>0.02</b> | 0.420 $\pm$ 0.03 | 0.520 $\pm$ 0.02                 |
| COAD Colorectal | 0.420 $\pm$ 0.06 | 0.670 $\pm$ 0.03 | 0.662 $\pm$ 0.01                 | 0.452 $\pm$ 0.02 | <b>0.729<math>\pm</math>0.03</b> |
| CCRCC Kidney    | 0.068 $\pm$ 0.01 | 0.055 $\pm$ 0.01 | 0.095 $\pm$ 0.01                 | 0.090 $\pm$ 0.01 | <b>0.109<math>\pm</math>0.01</b> |
| LUNG Lung       | 0.493 $\pm$ 0.07 | 0.378 $\pm$ 0.18 | 0.624 $\pm$ 0.01                 | 0.512 $\pm$ 0.03 | <b>0.675<math>\pm</math>0.01</b> |
| PAAD Pancreas   | 0.459 $\pm$ 0.02 | 0.399 $\pm$ 0.02 | 0.537 $\pm$ 0.02                 | 0.409 $\pm$ 0.03 | <b>0.548<math>\pm</math>0.01</b> |
| PRAD Prostate   | 0.079 $\pm$ 0.00 | 0.098 $\pm$ 0.02 | 0.118 $\pm$ 0.01                 | 0.046 $\pm$ 0.00 | <b>0.118<math>\pm</math>0.01</b> |
| SKCM Skin       | 0.520 $\pm$ 0.16 | 0.736 $\pm$ 0.01 | <b>0.748<math>\pm</math>0.01</b> | 0.631 $\pm$ 0.03 | 0.728 $\pm$ 0.04                 |
| Average         | 0.319            | 0.399            | 0.473                            | 0.366            | <b>0.490</b>                     |

STEM: corrected DiT ( $T = 1000$ , 34M params), three seeds.

Table 5: Cross-organ POOLED on HEST-1k. 5-fold pooled cross-validation; mean<sub>std</sub> over three seeds; nested-validation. Best per metric in bold.

| Method                                               | PCC $\uparrow$              | MSE $\downarrow$            | RMSE $\downarrow$           | MAE $\downarrow$            |
|------------------------------------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|
| ST-Net <a href="#">He et al. (2020)</a> (NBE'20)     | 0.370 <sub>0.01</sub>       | 0.944 <sub>0.05</sub>       | 0.965 <sub>0.02</sub>       | 0.653 <sub>0.01</sub>       |
| Hist2ST <a href="#">Zeng et al. (2022)</a> (BIB'22)  | 0.737 <sub>0.01</sub>       | 0.467 <sub>0.01</sub>       | 0.676 <sub>0.01</sub>       | 0.419 <sub>0.01</sub>       |
| BLEEP <a href="#">Xie et al. (2023)</a> (NeurIPS'23) | 0.758 <sub>0.01</sub>       | <b>0.427<sub>0.02</sub></b> | <b>0.648<sub>0.01</sub></b> | <b>0.380<sub>0.00</sub></b> |
| STEM <a href="#">Zhu et al. (2025)</a> (ICLR'25)     | 0.643 <sub>0.01</sub>       | 0.594 <sub>0.01</sub>       | 0.766 <sub>0.01</sub>       | 0.428 <sub>0.01</sub>       |
| MIST (V3)                                            | <b>0.760<sub>0.01</sub></b> | 0.440 <sub>0.02</sub>       | 0.657 <sub>0.02</sub>       | 0.410 <sub>0.01</sub>       |




Figure 2: Feature-encoder ablation (cross-organ POOLED, Pearson  $\uparrow$ ). Every method is trained identically under the POOLED regime (11-gene leakage-safe panel, three seeds) with three frozen patch encoders—UNI (1024-d), CONCH (512-d), and UNI + CONCH (1536-d). MIST (orange) attains the highest cross-organ transfer under every encoder, and UNI + CONCH is best overall (0.774; +0.014 over UNI, +0.056 over CONCH). Unlike the neighbourhood and context ablations, the encoder choice produces a real, consistent gap—the frozen features are the dominant factor.

Our panel comparison uses three 50-gene target sets shared by all evaluated models: HVG, the default HEST panel; DEG, selected by Wilcoxon tests across unsupervised Leiden clusters; and HMHVG, selected using expression mean and variability after variance-based candidate filtering. The models are trained and evaluated separately for each panel with the same cohort splits. Table 3 and Appendix G report the comparison.

Experiment Setup. Evaluation regimes. Intra-cohort evaluation trains a separate model within each cohort. Pooled five-fold cross-validation trains jointly across organs and holds out slide-level samples while retaining each organ type in training. Organ identity is used to define and report the cross-organ protocol, but organ labels and other slide metadata are not provided as model inputs. LOOO trains on six organs and tests the seventh, which is absent from training. Pooled folds are sample-disjoint; because patient identifiers are unavailable for some cohorts, we do not claim that they are patient-disjoint. To the best of our knowledge, this is the first pooled cross-organ evaluation of histology-to-spatial-expression prediction under a matched, controlled protocol.

Metrics and protocol. We report the Pearson correlation coefficient (PCC), Spearman correlation (SCC), and mean squared error (MSE) between predicted and measured expression, averaged over folds and three random seeds ( $\text{mean}_{\text{std}}$ ). Every model selects its checkpoint on an inner-validation split of the training fold and is evaluated once on the outer test-fold; where a cohort fold contains a single training slide, the inner split holds out a disjoint subset of that slide’s spots. This removes the test-fold model selection present in earlier exploratory runs.

Baselines. We compare with ST-Net [He et al. \(2020\)](#), Hist2ST [Zeng et al. \(2022\)](#), BLEEP [Xie et al. \(2023\)](#), and STEM [Zhu et al. \(2025\)](#) where results are available. Features, gene panels, splits, optimization budgets, and checkpoint selection are matched in the controlled tables. Published MCToGene results [Sun et al. \(2026\)](#), which use a different ten-cohort protocol, appear separately in Appendix F and are not treated as directly comparable.

Implementation. MIST uses  $L=4$  layers, hidden width  $d=256$ , 4 attention heads,  $k=8$  neighbours, and 16 radial-basis functions (4.8M parameters). We train per slide with AdamW (learning rate  $10^{-3}$ , weight decay  $10^{-2}$ ), a cosine schedule, and gradient clipping at 1.0, for at most 100 epochs with early-stopping patience of 20; each step samples at most 3,000 spots. Experiments run on one NVIDIA A100 GPU.

#### 4.1 Results for RQ1: Intra-cohort prediction

Table 2 reports the feature-matched comparison across gene-panel sizes. For PCC@50, MIST obtains the highest mean PCC (0.470), compared with 0.426 for BLEEP, the strongest baseline—a 10.3% relative improvement. It leads on all seven cohorts, though the margin is not uniform: it is widest on kidney and breast (+0.068 and +0.062 over the best baseline) and narrowest on colorectal (+0.004). The same ranking holds under a smaller 10-gene panel (PCC@10 in Table 2): with everything else fixed and only the panel reduced to the ten highest-variance genes of each cohort’s HEST 50-HVG set, MIST again leads the average (0.569 vs. 0.526 for BLEEP) and tops five of seven cohorts. Absolute correlations rise across all methods because the retained genes are the most variable, and hence the most predictable.

#### 4.2 Results for RQ2: Cross-organ transfer

LOOO is substantially harder because the target anatomy is absent from training. MIST reaches a mean PCC of 0.490 (Table 4), compared with 0.473 for BLEEP, the strongest baseline, and leads on five of seven held-out cohorts; BLEEP is stronger on IDC and SKCM. Under pooled cross-validation, where slide-level samples are held out while organ types remain represented in training, MIST reaches 0.760 (Table 5), narrowly above BLEEP at 0.758. The gap between pooled (0.760) and LOOO (0.490) directly quantifies the remaining difficulty of unseen-organ anatomical transfer. Appendix A reports a complementary seven-organ mouse experiment.

#### 4.3 Results for RQ3: What drives accuracy and coordinate invariance?

The factorial study in Table 6 evaluates all eight combinations of local (L), global (G), and slide-level (S) context on SKCM and LUNG. Local attention provides the largest gain:

Table 6: Factorial context ablation (intra-cohort, Pearson mean<sub>std</sub> over three seeds on SKCM and LUNG). L, G, and S denote local attention, global attention, and slide representation. Local attention is the dominant contributor within an organ.

| Variant | L | G | S | SKCM                        | LUNG                        | Avg          |
|---------|---|---|---|-----------------------------|-----------------------------|--------------|
| C000    |   |   |   | 0.639 <sub>0.00</sub>       | 0.577 <sub>0.01</sub>       | 0.608        |
| C001    |   |   | • | 0.633 <sub>0.01</sub>       | 0.578 <sub>0.00</sub>       | 0.605        |
| C010    |   | • |   | 0.641 <sub>0.01</sub>       | 0.577 <sub>0.00</sub>       | 0.609        |
| C011    |   | • | • | 0.639 <sub>0.02</sub>       | 0.583 <sub>0.00</sub>       | 0.611        |
| C100    | • |   |   | 0.693 <sub>0.01</sub>       | 0.587 <sub>0.00</sub>       | 0.640        |
| C101    | • |   | • | 0.695 <sub>0.01</sub>       | 0.585 <sub>0.01</sub>       | 0.640        |
| C110    | • | • |   | 0.693 <sub>0.01</sub>       | 0.584 <sub>0.00</sub>       | 0.638        |
| C111    | • | • | • | <b>0.696<sub>0.00</sub></b> | <b>0.589<sub>0.01</sub></b> | <b>0.643</b> |

Table 7: Numerical coordinate-invariance check. Maximum absolute prediction difference when only the spot coordinates are transformed (randomly initialized model, 257-spot synthetic slide).

| Transformation   | Max. abs. difference  |
|------------------|-----------------------|
| Translation      | $5.97 \times 10^{-5}$ |
| Rotation         | $1.97 \times 10^{-5}$ |
| Reflection       | 0                     |
| Spot permutation | $5.96 \times 10^{-7}$ |

enabling it alone raises mean PCC from 0.608 to 0.640. Global and slide context have little effect without the local pathway, but the complete model achieves the highest mean (0.643). This result supports distance-biased neighbourhood aggregation as the primary contributor in the tested intra-cohort setting; claims about component importance under organ shift require a matched LOOO ablation not reported here.

Table 7 numerically verifies the analytic invariance result. Differences under translation and rotation remain below  $6 \times 10^{-5}$  and are attributable to floating-point evaluation; reflection is exact in this test.




Figure 3: Qualitative spatial expression maps for a held-out skin (SKCM) slide (TENX117), three genes (rows). Each panel colours every spot by expression; MIST (right) most faithfully recovers the ground-truth spatial structure—the perivascular ring for LYZ/CD52 and the epidermal edge for KRTDAP—attaining the highest per-gene PCC (in orange), whereas STEM is noisy and Hist2ST over-smooths. Predictions use identical frozen UNI+CONCH features.

#### 4.4 Results for RQ4: Spatial structure preservation

Gene-wise PCC does not test whether predictions preserve coherent tissue regions. We cluster spots using either predicted or measured expression and compare the induced domains with adjusted Rand index (ARI) and NMI. MIST leads on both ARI (0.184) and NMI (0.239), ahead of the strongest baseline (ST-Net at 0.154 and 0.212; Appendix B). The differences remain small and slide-level variation is large, so we caution against over-interpreting them as evidence of uniformly improved domain recovery.

#### 4.5 Results for RQ5: Computational efficiency

MIST contains 4.88M parameters and predicts a complete slide in one forward pass (Table 8). Its advantage over the baselines is not minimum parameter count — ST-Net and BLEEP are smaller — but the combination of multi-scale spatial context and direct inference. Model-only latency grows from 3.5 ms for 500 spots to 64.7 ms for 10,000 spots on an A100, with 3.1 GiB peak allocated memory at the largest size (Table 9); feature extraction and disk I/O are excluded. STEM’s iterative diffusion inference is the only approach with substantially higher latency.

Table 8: Model complexity. Parameter counts and inference mechanisms.

| Model                                                | #Params (M) | Inference   |
|------------------------------------------------------|-------------|-------------|
| ST-Net <a href="#">He et al. (2020)</a> (NBE’20)     | 0.47        | single-pass |
| Hist2ST <a href="#">Zeng et al. (2022)</a> (BIB’22)  | 2.18        | single-pass |
| BLEEP <a href="#">Xie et al. (2023)</a> (NeurIPS’23) | 0.54        | single-pass |
| STEM <a href="#">Zhu et al. (2025)</a> (ICLR’25)     | 34.16       | iterative   |
| MIST (V3)                                            | 4.88        | single-pass |

Table 9: MIST inference scaling on an NVIDIA A100 80GB. Median model-only latency and peak allocated GPU memory over synthetic slides; feature extraction and disk I/O are excluded.

| Spots  | Latency (ms) | Memory (MiB) |
|--------|--------------|--------------|
| 500    | 3.48         | 43           |
| 1,000  | 3.63         | 67           |
| 3,000  | 9.36         | 335          |
| 5,000  | 21.31        | 842          |
| 10,000 | 64.74        | 3,183        |

## 5 Discussion

Three findings merit emphasis. First, a compact direct predictor is competitive when features, gene panels, splits, and model selection are controlled; the 10.3% relative gain over the strongest baseline rests on matched experimental conditions, not a larger model. Second, pooled multi-organ performance is not a substitute for unseen-organ evaluation: the drop from 0.760 (pooled) to 0.490 (LOOO) shows that anatomical shift remains a separate and harder problem. Third, local distance-biased attention is the dominant contributor in the factorial study; global and slide pathways provide moderate complementary gains, and the complete model (C111) achieves the best mean across both tested cohorts.

Direct and generative models serve different purposes. STEM [Zhu et al. \(2025\)](#) models a conditional expression distribution and can represent one-to-many uncertainty through iterative sampling. MIST produces a point estimate optimised for correlation and squared error. Its benefits — deterministic, single-pass inference and explicit coordinate invariance — do not imply that morphology uniquely determines expression.

Limitations. The evaluation predicts 50 genes from frozen UNI+CONCH embeddings; results may not extend to transcriptome-scale output or end-to-end image encoding. LOOO uses a fixed benchmark panel and does not test platform or institution shift simultaneously. Coordinate invariance holds only for the spatial pathway with fixed image features; rotating the histology image can change its frozen embedding. Domain-recovery analysis covers only two baselines and has high slide-level variance. Pooled folds are sample-disjoint but cannot be verified as patient-disjoint where patient identifiers are unavailable.

## 6 Conclusion and Future Work

We developed MIST, a coordinate-invariant local-global spatial transformer for predicting gene expression from histology. Pairwise-distance encoding makes the spatial pathway invariant to rigid coordinate transformations by construction, while three parallel context signals — local  $k$ NN attention, global self-attention, and slide-level pooling — together produce complete-slide predictions in a single pass. These design choices directly address our design goals that multi-scale context is necessary and that architectural positional bias should not break under coordinate shift. Under matched HEST-1k evaluation, MIST improves mean intra-cohort correlation and remains competitive in LOOO transfer; the pooled-to-LOOO gap exposes the remaining difficulty of anatomical transfer as the central open challenge.

For future work, it would be valuable to explore how conditioning on slide-level metadata (organ type, sequencing platform) can narrow this anatomical gap — for instance, through FiLM or adaptive layer normalisation conditioners that are available at inference time without requiring expression labels. Scaling the gene panel beyond 50 genes, incorporating end-to-end image encoding, and combining MIST’s geometric structure with calibrated generative uncertainty modelling (as in STEM) are further natural directions.

#### AI use statement

Generative AI tools assisted with manuscript organization, language editing, ICLR formatting, and consistency checks between narrative claims and reported tables. They were not used to generate experimental data or execute model experiments. The authors reviewed all AI-assisted revisions and take responsibility for the final content of this work.

#### References

- Richard J Chen, Tong Ding, Ming Y Lu, et al. Towards a general-purpose foundation model for computational pathology. *Nature Medicine*, 30:850–862, 2024. 4
- Bryan He, Ludvig Bergenstråhle, Linnea Stenbeck, et al. Integrating spatial gene expression and breast tumour morphology via deep learning. *Nature Biomedical Engineering*, 4(8): 827–834, 2020. 1, 3, 6, 9, 10, 13, 15, 17, 18, 20, 21
- Guillaume Jaume, Paul Doucet, Andrew H Song, et al. HEST-1k: A dataset for spatial transcriptomics and histology image analysis. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 4, 6
- Ming Y. Lu, Bowen Chen, Drew F. K. Williamson, et al. A visual-language foundation model for computational pathology. *Nature Medicine*, 30:863–874, 2024. doi: 10.1038/s41591-024-02856-4. 4
- Victor Garcia Satorras, Emiel Hoogetboom, and Max Welling. E(n) equivariant graph neural networks. In *International Conference on Machine Learning (ICML)*, 2021. 4
- Patrik L Ståhl, Fredrik Salmén, Sanja Vickovic, et al. Visualization and analysis of gene expression in tissue sections by spatial transcriptomics. *Science*, 353(6294):78–82, 2016. 1
- Youhan Sun, Jiahua Rao, Kangrui Du, Jiancong Xie, and Yuedong Yang. Predicting spatial transcriptomics from histology images via high-order multi-cell interaction modeling. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. 10, 17
- Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. 5
- Ronald Xie, Kuan Pang, Sai Chung, et al. Spatially resolved gene expression prediction from histology images via bi-modal contrastive learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. 1, 3, 4, 6, 9, 10, 13, 15, 18, 20, 21
- Yuansong Zeng, Zhuoyi Wei, Weijiang Yu, et al. Spatial transcriptomics prediction from histology jointly through transformer and graph neural networks. *Briefings in Bioinformatics*, 23(5), 2022. 1, 2, 3, 6, 9, 10, 13, 15, 17, 18, 20, 21
- Sichen Zhu, Yuchen Zhu, Molei Tao, and Peng Qiu. Diffusion generative modeling for spatially resolved gene expression inference from histology images. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=FtjLUHyZAO>. 1, 3, 4, 6, 9, 10, 13, 15, 18, 20, 21

## A Mouse cross-organ diagnostic

We evaluate 59 mouse Visium slides spanning seven organs with mouse-specific 50-gene panels. Pooled performance reaches 0.887, whereas every method remains near the floor under LOOO; we therefore do not interpret the latter as evidence of successful cross-species organ transfer.

Table 10: Cross-organ generalization on mouse HEST (three seeds,  $\text{mean}_{\text{std}}$ ). looo: leave-one-organ-out transfer; pooled: five-fold pooled cross-validation. PCC $\uparrow$  on 50 mouse HVGs.

| Method                                     | looo PCC $\uparrow$          | pooled PCC                   |
|--------------------------------------------|------------------------------|------------------------------|
| ST-Net <a href="#">He et al. (2020)</a>    | 0.013 <sub>0.011</sub>       | 0.357 <sub>0.036</sub>       |
| Hist2ST <a href="#">Zeng et al. (2022)</a> | 0.018 <sub>0.007</sub>       | 0.829 <sub>0.028</sub>       |
| BLEEP <a href="#">Xie et al. (2023)</a>    | -0.006 <sub>0.003</sub>      | 0.848 <sub>0.012</sub>       |
| STEM <a href="#">Zhu et al. (2025)</a>     | —                            | 0.279 <sub>0.009</sub>       |
| MIST                                       | <b>0.018<sub>0.006</sub></b> | <b>0.887<sub>0.005</sub></b> |

$\dagger$ All looo values are near the floor because the organ-specific gene panels do not transfer reliably.

## B Spatial-domain recovery

Table 11: Spatial-domain recovery. ARI/NMI between domains clustered from predicted and measured expression on held-out slides ( $\text{mean}_{\text{std}}$  over three seeds; higher is better).

| Method                                     | ARI $\uparrow$              | NMI $\uparrow$              |
|--------------------------------------------|-----------------------------|-----------------------------|
| ST-Net <a href="#">He et al. (2020)</a>    | 0.154 <sub>0.14</sub>       | 0.212 <sub>0.15</sub>       |
| Hist2ST <a href="#">Zeng et al. (2022)</a> | 0.134 <sub>0.11</sub>       | 0.180 <sub>0.11</sub>       |
| BLEEP <a href="#">Xie et al. (2023)</a>    | 0.147 <sub>0.12</sub>       | 0.186 <sub>0.13</sub>       |
| STEM <a href="#">Zhu et al. (2025)</a>     | 0.109 <sub>0.12</sub>       | 0.131 <sub>0.13</sub>       |
| MIST                                       | <b>0.184<sub>0.16</sub></b> | <b>0.239<sub>0.16</sub></b> |

## C Full per-cohort results

Table 12 gives the per-cohort breakdown for all seven retained cohorts (HCC/liver, LYMPH\_IDC, and READ excluded; see Appendix H) across all six baselines and MIST.

## D Cross-organ LOOO transfer, expanded baselines (UNI + CONCH)

Table 13 extends Table 4 with M2OST and EGN under the same UNI+CONCH protocol. EGN attains the highest mean PCC (0.500), followed by MIST at 0.490; their per-cohort ranking varies.

## E Cross-organ LOOO transfer with UNI features

Table 14 reports a complementary LOOO comparison using UNI features. EGN has the highest mean PCC (0.493), while MIST reaches 0.477. M2OST narrowly leads EGN on LUNG (0.667 versus 0.666). STEM uses the corrected DiT configuration with  $T = 1000$ ; its SKCM entry contains one completed seed and is marked accordingly.

Table 12: Full per-cohort intra-cohort results on HEST-1k (Pearson  $\text{mean}_{\text{std}}$  over three seeds; nested-validation). All methods use identical frozen UNI + CONCH features (1536-dim) and 50-gene panels, except MIST (V3) which uses UNI (1024-dim) for comparison. HCC (liver), LYMPH\_IDC, and READ are excluded per Appendix H. Best per row in bold.

| Cohort          | ST-Net NBE'20    | Hist2ST BIB'22   | BLEEP NeurIPS'23 | M2OST AAAI'25                    | EGN WACV'23      | STEM ICLR'25     | MIST (V3)                        | MIST+CONCH                       |
|-----------------|------------------|------------------|------------------|----------------------------------|------------------|------------------|----------------------------------|----------------------------------|
| IDC Breast      | 0.551 $\pm$ 0.01 | 0.309 $\pm$ 0.07 | 0.519 $\pm$ 0.00 | 0.571 $\pm$ 0.01                 | 0.551 $\pm$ 0.01 | 0.439 $\pm$ 0.01 | 0.567 $\pm$ 0.00                 | <b>0.581<math>\pm</math>0.01</b> |
| COAD Colorectal | 0.242 $\pm$ 0.01 | 0.295 $\pm$ 0.01 | 0.314 $\pm$ 0.01 | <b>0.328<math>\pm</math>0.01</b> | 0.293 $\pm$ 0.01 | 0.177 $\pm$ 0.01 | 0.299 $\pm$ 0.00                 | 0.318 $\pm$ 0.01                 |
| CCRCC Kidney    | 0.163 $\pm$ 0.02 | 0.184 $\pm$ 0.02 | 0.175 $\pm$ 0.02 | 0.210 $\pm$ 0.02                 | 0.223 $\pm$ 0.02 | 0.123 $\pm$ 0.00 | <b>0.265<math>\pm</math>0.02</b> | 0.252 $\pm$ 0.01                 |
| LUNG Lung       | 0.518 $\pm$ 0.01 | 0.470 $\pm$ 0.01 | 0.562 $\pm$ 0.01 | <b>0.598<math>\pm</math>0.00</b> | 0.584 $\pm$ 0.01 | 0.423 $\pm$ 0.03 | 0.589 $\pm$ 0.01                 | 0.589 $\pm$ 0.01                 |
| PAAD Pancreas   | 0.359 $\pm$ 0.03 | 0.247 $\pm$ 0.02 | 0.419 $\pm$ 0.03 | 0.413 $\pm$ 0.04                 | 0.415 $\pm$ 0.03 | 0.233 $\pm$ 0.02 | 0.392 $\pm$ 0.04                 | <b>0.437<math>\pm</math>0.03</b> |
| PRAD Prostate   | 0.294 $\pm$ 0.02 | 0.333 $\pm$ 0.02 | 0.348 $\pm$ 0.01 | <b>0.413<math>\pm</math>0.01</b> | 0.377 $\pm$ 0.03 | 0.207 $\pm$ 0.02 | 0.390 $\pm$ 0.02                 | 0.410 $\pm$ 0.01                 |
| SKCM Skin       | 0.665 $\pm$ 0.01 | 0.522 $\pm$ 0.03 | 0.645 $\pm$ 0.01 | 0.654 $\pm$ 0.01                 | 0.631 $\pm$ 0.01 | 0.491 $\pm$ 0.04 | 0.696 $\pm$ 0.00                 | <b>0.707<math>\pm</math>0.01</b> |
| Average         | 0.399            | 0.337            | 0.426            | 0.455                            | 0.439            | 0.299            | 0.457                            | <b>0.470</b>                     |

MIST (V3) uses UNI (1024-dim); all other methods use UNI + CONCH (1536-dim). STEM: corrected DiT (T=1000, 34M params).

Table 13: Cross-organ LOOO transfer, expanded baselines, per-cohort Pearson  $\uparrow$  (HEST-1k) (Pearson  $\text{mean}_{\text{std}}$  over three seeds; nested-validation; liver excluded). Leave-one-organ-out: trained on all other organ groups, tested on the held-out cohort. All methods use identical frozen UNI + CONCH features (1536-dim) and 50-gene panels. Best per row in bold.

| Cohort          | ST-Net NBE'20    | Hist2ST BIB'22   | BLEEP NeurIPS'23 | M2OST AAAI'25    | EGN WACV'23                      | STEM ICLR'25     | MIST+CONCH                       |
|-----------------|------------------|------------------|------------------|------------------|----------------------------------|------------------|----------------------------------|
| IDC Breast      | 0.198 $\pm$ 0.05 | 0.458 $\pm$ 0.01 | 0.529 $\pm$ 0.02 | 0.437 $\pm$ 0.04 | <b>0.567<math>\pm</math>0.02</b> | 0.420 $\pm$ 0.03 | 0.520 $\pm$ 0.02                 |
| COAD Colorectal | 0.420 $\pm$ 0.06 | 0.670 $\pm$ 0.03 | 0.662 $\pm$ 0.01 | 0.657 $\pm$ 0.05 | 0.728 $\pm$ 0.01                 | 0.452 $\pm$ 0.02 | <b>0.729<math>\pm</math>0.03</b> |
| CCRCC Kidney    | 0.068 $\pm$ 0.01 | 0.055 $\pm$ 0.01 | 0.095 $\pm$ 0.01 | 0.081 $\pm$ 0.03 | 0.103 $\pm$ 0.02                 | 0.090 $\pm$ 0.01 | <b>0.109<math>\pm</math>0.01</b> |
| LUNG Lung       | 0.493 $\pm$ 0.07 | 0.378 $\pm$ 0.18 | 0.624 $\pm$ 0.01 | 0.665 $\pm$ 0.01 | 0.672 $\pm$ 0.01                 | 0.512 $\pm$ 0.03 | <b>0.675<math>\pm</math>0.01</b> |
| PAAD Pancreas   | 0.459 $\pm$ 0.02 | 0.399 $\pm$ 0.02 | 0.537 $\pm$ 0.02 | 0.513 $\pm$ 0.03 | <b>0.572<math>\pm</math>0.01</b> | 0.409 $\pm$ 0.03 | 0.548 $\pm$ 0.01                 |
| PRAD Prostate   | 0.079 $\pm$ 0.00 | 0.098 $\pm$ 0.02 | 0.118 $\pm$ 0.01 | 0.087 $\pm$ 0.02 | 0.107 $\pm$ 0.01                 | 0.046 $\pm$ 0.00 | <b>0.118<math>\pm</math>0.01</b> |
| SKCM Skin       | 0.520 $\pm$ 0.16 | 0.736 $\pm$ 0.01 | 0.748 $\pm$ 0.01 | 0.740 $\pm$ 0.01 | <b>0.750<math>\pm</math>0.02</b> | 0.631 $\pm$ 0.03 | 0.728 $\pm$ 0.04                 |
| Average         | 0.319            | 0.399            | 0.473            | 0.454            | <b>0.500</b>                     | 0.366            | 0.490                            |

## F Published Per-Cohort Comparison

## G Gene Panel Comparison

Table 16 reports intra-cohort mean PCC averaged over seven cohorts and three seeds for three gene-panel strategies: highly variable genes (HVG, the HEST-1k default), unsupervised differentially expressed genes (DEG; Leiden clustering + Wilcoxon), and high-mean highly variable genes (HMHVG; selected using expression mean and variability after variance filtering). The supplied panel-generation script constructs DEG and HMHVG panels once per cohort using all available slides, rather than selecting genes separately within each training fold. These results therefore describe an exploratory comparison with fixed cohort-level panels; held-out expression can influence target selection, and nested check-point validation does not remove this dependence. A strictly inductive comparison requires training-fold-only panel selection and retraining. All methods use identical UNI+CONCH features, nested-validation, and the same training protocol. DEG and HMHVG panels consistently outperform HVG across all methods, with HMHVG yielding the highest mean for every method tested. MIST leads on all three panels, with its advantage most pronounced on HMHVG (0.657 vs. 0.595 for BLEEP) and consistent across DEG (0.624 vs. 0.571) and HVG (0.471 vs. 0.427).

Table 17 gives the per-cohort breakdown. MIST leads on 6 of 7 cohorts under HVG (BLEEP is stronger on COAD) and on all 7 cohorts under both DEG and HMHVG.

## H HEST-1k Cohort Statistics

Table 18 lists the number of samples, total spots, and cross-validation folds for each HEST-1k cohort. Two organ groups in HEST-1k contain multiple cohorts: breast (IDC + LYMPH\_IDC) and colorectal (COAD + READ). Within each group the constituent

Table 14: Cross-organ LOOO transfer with UNI features (Pearson  $\text{mean}_{\text{std}}$  over three seeds; nested-validation; HCC, LYMPH\_IDC, and READ excluded). Leave-one-organ-out: trained on all other organs, tested on the held-out organ. Best per row in bold.

| Organ           | ST-Net NBE'20 | Hist2ST BIB'22 | BLEEP NeurIPS'23 | M2OST AAAI'25     | EGN WACV'23       | STEM ICLR'25       | MIST (V3)         |
|-----------------|---------------|----------------|------------------|-------------------|-------------------|--------------------|-------------------|
| IDC Breast      | 0.217±0.06    | 0.459±0.02     | 0.509±0.03       | 0.452±0.01        | 0.531±0.03        | 0.406±0.01         | <b>0.534±0.02</b> |
| COAD Colorectal | 0.430±0.03    | 0.688±0.04     | 0.667±0.02       | 0.672±0.03        | <b>0.749±0.01</b> | 0.504±0.03         | 0.703±0.01        |
| CCRCC Kidney    | 0.084±0.03    | 0.084±0.01     | 0.100±0.02       | 0.122±0.01        | 0.106±0.01        | 0.090±0.01         | <b>0.124±0.00</b> |
| LUNG Lung       | 0.504±0.04    | 0.499±0.06     | 0.621±0.02       | <b>0.667±0.01</b> | 0.666±0.01        | 0.561±0.01         | 0.660±0.01        |
| PAAD Pancreas   | 0.404±0.05    | 0.360±0.07     | 0.521±0.02       | 0.522±0.01        | <b>0.551±0.02</b> | 0.460±0.02         | 0.507±0.04        |
| PRAD Prostate   | 0.071±0.01    | 0.104±0.02     | 0.111±0.00       | 0.107±0.02        | 0.098±0.02        | 0.068±0.01         | <b>0.119±0.02</b> |
| SKCM Skin       | 0.402±0.08    | 0.652±0.08     | 0.745±0.01       | 0.724±0.02        | <b>0.749±0.01</b> | 0.686 <sup>†</sup> | 0.692±0.06        |
| Average         | 0.302         | 0.407          | 0.468            | 0.466             | <b>0.493</b>      | 0.396 <sup>†</sup> | 0.477             |

<sup>†</sup>STEM DiT (corrected, T=1000, 100 epochs) LOOO: SKCM (skin) has only 1 seed complete; all other organs use 3 seeds. Average marked <sup>†</sup> accordingly.

Table 15: Published per-cohort accuracy on the standard HEST-1k splits, mean PCC<sup>†</sup> as reported in respective papers (ten cohorts, no organ grouping). Shown for context only—not comparable to Table 2, which uses organ grouping and nested-validation. <sup>†</sup>MIST uses our organ-grouped splits, not the standard per-cohort split.

| Method                               | HEST-1k (10 cohorts) |
|--------------------------------------|----------------------|
| ST-Net He et al. (2020) (NBE'20)     | 0.286                |
| Hist2ST Zeng et al. (2022) (BIB'22)  | —                    |
| MCToGene Sun et al. (2026) (CVPR'26) | <b>0.435</b>         |
| MIST <sup>†</sup> (ours)             | 0.440                |

cohorts differ substantially in size: IDC has  $2.25\times$  more spots than LYMPH\_IDC (44,974 vs. 19,964), and COAD has  $2.2\times$  more spots than READ (18,523 vs. 8,407). To avoid misrepresentation from unweighted averaging over unequal cohorts, Table 2 uses IDC as the representative breast cohort and COAD as the representative colorectal cohort, discarding the smaller partner cohorts (LYMPH\_IDC and READ) from the organ-level rows. This choice maximises statistical reliability of each organ-level estimate.

Table 16: Gene panel comparison: intra-cohort mean PCC $\uparrow$  (HEST-1k). Mean<sub>std</sub> over three seeds; per-seed value is the average over seven cohorts. All methods use identical frozen UNI + CONCH features (1536-dim) and nested-validation. Best per column in bold.

| Method                                     | HVG                          | DEG                          | HMHVG                        |
|--------------------------------------------|------------------------------|------------------------------|------------------------------|
| ST-Net <a href="#">He et al. (2020)</a>    | 0.398 <sub>0.004</sub>       | 0.523 <sub>0.004</sub>       | 0.550 <sub>0.008</sub>       |
| Hist2ST <a href="#">Zeng et al. (2022)</a> | 0.330 <sub>0.032</sub>       | 0.494 <sub>0.014</sub>       | 0.544 <sub>0.038</sub>       |
| BLEEP <a href="#">Xie et al. (2023)</a>    | 0.427 <sub>0.004</sub>       | 0.571 <sub>0.005</sub>       | 0.595 <sub>0.014</sub>       |
| STEM <a href="#">Zhu et al. (2025)</a>     | 0.291 <sub>0.005</sub>       | 0.456 <sub>0.002</sub>       | 0.485 <sub>0.010</sub>       |
| MIST (ours)                                | <b>0.471<sub>0.005</sub></b> | <b>0.624<sub>0.002</sub></b> | <b>0.657<sub>0.006</sub></b> |

Table 17: Per-cohort gene panel comparison: intra-cohort PCC $\uparrow$  (HEST-1k). Mean<sub>std</sub> over three seeds. All methods use UNI + CONCH (1536-dim) and nested-validation. Best per cohort per panel in bold.

| Panel | Cohort          | ST-Net                 | Hist2ST                | BLEEP                        | STEM                   | MIST                         |
|-------|-----------------|------------------------|------------------------|------------------------------|------------------------|------------------------------|
| HVG   | IDC Breast      | 0.551 <sub>0.007</sub> | 0.305 <sub>0.108</sub> | 0.519 <sub>0.002</sub>       | 0.434 <sub>0.002</sub> | <b>0.581<sub>0.008</sub></b> |
|       | COAD Colorectal | 0.235 <sub>0.005</sub> | 0.298 <sub>0.009</sub> | <b>0.322<sub>0.006</sub></b> | 0.177 <sub>0.014</sub> | 0.318 <sub>0.008</sub>       |
|       | CCRCC Kidney    | 0.163 <sub>0.015</sub> | 0.180 <sub>0.030</sub> | 0.175 <sub>0.015</sub>       | 0.118 <sub>0.006</sub> | <b>0.252<sub>0.007</sub></b> |
|       | LUNG Lung       | 0.521 <sub>0.010</sub> | 0.430 <sub>0.050</sub> | 0.565 <sub>0.008</sub>       | 0.407 <sub>0.003</sub> | <b>0.589<sub>0.006</sub></b> |
|       | PAAD Pancreas   | 0.366 <sub>0.032</sub> | 0.248 <sub>0.038</sub> | 0.412 <sub>0.034</sub>       | 0.233 <sub>0.019</sub> | <b>0.437<sub>0.028</sub></b> |
|       | PRAD Prostate   | 0.287 <sub>0.023</sub> | 0.323 <sub>0.012</sub> | 0.346 <sub>0.007</sub>       | 0.207 <sub>0.016</sub> | <b>0.410<sub>0.006</sub></b> |
|       | SKCM Skin       | 0.661 <sub>0.006</sub> | 0.523 <sub>0.003</sub> | 0.650 <sub>0.004</sub>       | 0.461 <sub>0.018</sub> | <b>0.707<sub>0.007</sub></b> |
|       | Average         | 0.398                  | 0.330                  | 0.427                        | 0.291                  | <b>0.471</b>                 |
| DEG   | IDC Breast      | 0.697 <sub>0.010</sub> | 0.549 <sub>0.067</sub> | 0.731 <sub>0.014</sub>       | 0.684 <sub>0.020</sub> | <b>0.774<sub>0.009</sub></b> |
|       | COAD Colorectal | 0.497 <sub>0.003</sub> | 0.608 <sub>0.010</sub> | 0.639 <sub>0.016</sub>       | 0.534 <sub>0.014</sub> | <b>0.677<sub>0.003</sub></b> |
|       | CCRCC Kidney    | 0.312 <sub>0.014</sub> | 0.335 <sub>0.013</sub> | 0.304 <sub>0.003</sub>       | 0.253 <sub>0.009</sub> | <b>0.402<sub>0.013</sub></b> |
|       | LUNG Lung       | 0.566 <sub>0.006</sub> | 0.518 <sub>0.006</sub> | 0.621 <sub>0.006</sub>       | 0.440 <sub>0.009</sub> | <b>0.662<sub>0.002</sub></b> |
|       | PAAD Pancreas   | 0.425 <sub>0.032</sub> | 0.352 <sub>0.011</sub> | 0.487 <sub>0.029</sub>       | 0.306 <sub>0.042</sub> | <b>0.518<sub>0.026</sub></b> |
|       | PRAD Prostate   | 0.413 <sub>0.004</sub> | 0.437 <sub>0.008</sub> | 0.466 <sub>0.005</sub>       | 0.355 <sub>0.026</sub> | <b>0.539<sub>0.012</sub></b> |
|       | SKCM Skin       | 0.752 <sub>0.009</sub> | 0.656 <sub>0.004</sub> | 0.746 <sub>0.026</sub>       | 0.617 <sub>0.023</sub> | <b>0.794<sub>0.002</sub></b> |
|       | Average         | 0.523                  | 0.494                  | 0.571                        | 0.456                  | <b>0.624</b>                 |
| HMHVG | IDC Breast      | 0.724 <sub>0.014</sub> | 0.633 <sub>0.082</sub> | 0.739 <sub>0.018</sub>       | 0.687 <sub>0.008</sub> | <b>0.774<sub>0.011</sub></b> |
|       | COAD Colorectal | 0.503 <sub>0.010</sub> | 0.659 <sub>0.018</sub> | 0.645 <sub>0.004</sub>       | 0.528 <sub>0.032</sub> | <b>0.698<sub>0.005</sub></b> |
|       | CCRCC Kidney    | 0.339 <sub>0.007</sub> | 0.387 <sub>0.010</sub> | 0.349 <sub>0.010</sub>       | 0.290 <sub>0.013</sub> | <b>0.454<sub>0.017</sub></b> |
|       | LUNG Lung       | 0.568 <sub>0.005</sub> | 0.598 <sub>0.003</sub> | 0.661 <sub>0.002</sub>       | 0.517 <sub>0.011</sub> | <b>0.693<sub>0.004</sub></b> |
|       | PAAD Pancreas   | 0.441 <sub>0.032</sub> | 0.341 <sub>0.038</sub> | 0.493 <sub>0.023</sub>       | 0.305 <sub>0.027</sub> | <b>0.524<sub>0.023</sub></b> |
|       | PRAD Prostate   | 0.452 <sub>0.002</sub> | 0.477 <sub>0.011</sub> | 0.432 <sub>0.045</sub>       | 0.337 <sub>0.016</sub> | <b>0.555<sub>0.022</sub></b> |
|       | SKCM Skin       | 0.824 <sub>0.029</sub> | 0.710 <sub>0.141</sub> | 0.849 <sub>0.018</sub>       | 0.731 <sub>0.060</sub> | <b>0.901<sub>0.002</sub></b> |
|       | Average         | 0.550                  | 0.544                  | 0.595                        | 0.485                  | <b>0.657</b>                 |

Table 18: HEST-1k cohort statistics. Spots counted after barcode filtering. Folds = number of train/test splits in the per-cohort cross-validation. HCC (liver) is excluded from Table 2 due to its near-zero signal floor.

| Cohort    | Organ      | Samples | Spots  | Folds |
|-----------|------------|---------|--------|-------|
| IDC       | Breast     | 4       | 44,974 | 4     |
| LYMPH_IDC | Breast     | 4       | 19,964 | 4     |
| COAD      | Colorectal | 4       | 18,523 | 2     |
| READ      | Colorectal | 4       | 8,407  | 2     |
| CCRCC     | Kidney     | 24      | 74,220 | 6     |
| HCC       | Liver*     | 2       | 4,248  | 2     |
| LUNG      | Lung       | 2       | 7,505  | 2     |
| PAAD      | Pancreas   | 3       | 9,793  | 3     |
| PRAD      | Prostate   | 23      | 62,710 | 2     |
| SKCM      | Skin       | 2       | 5,716  | 2     |

\*HCC (liver) is included here for completeness but excluded from Table 2; all baselines score near zero on this cohort due to low spatial gene-expression variance.

## I Proofs of the formal guarantees

### I.1 Proof of Proposition 1

For every pair of spots  $i, j$ , orthogonality of  $R$  gives

$$\|R\mathbf{p}_i + \mathbf{t} - (R\mathbf{p}_j + \mathbf{t})\|_2 = \|R(\mathbf{p}_i - \mathbf{p}_j)\|_2 \quad (9)$$

$$= \|\mathbf{p}_i - \mathbf{p}_j\|_2. \quad (10)$$

Hence, transforming the coordinate frame leaves the complete distance matrix unchanged. It also leaves the  $k$ NN graph, radial basis features, and distance-dependent attention biases unchanged. The input projection is unchanged because  $\mathbf{F}$  is fixed. At every layer, the local pathway therefore receives the same features, neighbours, and distance biases, while the global and slide pathways do not use coordinates. Their outputs, the residual fusion, and the feed-forward update are identical. Induction over the  $L$  layers gives identical final representations and predictions, proving Proposition 1.

### I.2 Proof of Proposition 2

Let  $\mathbf{F}' = Q\mathbf{F}$  and  $\mathbf{P}' = Q\mathbf{P}$ . The pairwise distance matrix is relabelled as  $\mathbf{D}' = Q\mathbf{D}Q^\top$ . Under equivariant tie handling, the  $k$ NN graph and its radial basis edge features are relabelled by the same permutation. The input projection and all pointwise linear maps, layer normalizations, activations, and feed-forward networks commute with  $Q$ .

For local attention, permuting the spots only relabels each neighbourhood and its attention sum, so the local output is multiplied by  $Q$ . Global self-attention is likewise permutation equivariant because its query-key score matrix is transformed by  $Q$  on both axes and its value matrix by  $Q$ . The slide attention scores are permuted, but their weighted sum is unchanged; broadcasting that same summary to all spots is equivariant. Consequently, residual fusion preserves equivariance at every layer. Induction over the layers and application of the pointwise prediction head yield  $f_\theta(Q\mathbf{F}, Q\mathbf{P}) = Qf_\theta(\mathbf{F}, \mathbf{P})$ , proving Proposition 2.

## J Three-Metric Intra-Cohort Summary

Table 19 extends the intra-cohort comparison with Spearman correlation (SCC) and mean squared error (MSE) alongside PCC, averaged over seven organ groups.

Table 19: Intra-cohort three-metric summary (HEST-1k, averaged over seven organ groups; three seeds; UNI + CONCH features). PCC $\uparrow$ , SCC $\uparrow$ , MSE $\downarrow$ . Best per column in bold.

| Method                                               | PCC $\uparrow$ | SCC $\uparrow$ | MSE $\downarrow$ |
|------------------------------------------------------|----------------|----------------|------------------|
| ST-Net <a href="#">He et al. (2020)</a> (NBE'20)     | 0.399          | 0.356          | 2.043            |
| Hist2ST <a href="#">Zeng et al. (2022)</a> (BIB'22)  | 0.337          | 0.321          | 1.616            |
| BLEEP <a href="#">Xie et al. (2023)</a> (NeurIPS'23) | 0.426          | 0.388          | 1.464            |
| STEM <a href="#">Zhu et al. (2025)</a> (ICLR'25)     | 0.299          | 0.274          | 1.855            |
| MIST (ours)                                          | 0.470          | 0.430          | 1.389            |

## K Per-Cohort MSE and MAE (Intra-Cohort)

Tables 20 and 21 report the per-cohort mean squared error (MSE) and mean absolute error (MAE) for the same models, cohorts, and protocol as the intra-cohort Pearson comparison in Table 2. MIST attains the best average on both error metrics (MSE 1.389, MAE 0.829), consistent with its Pearson result; per-cohort leadership is less uniform than on Pearson, as the training objective couples MSE with a correlation term.

Table 20: Intra-cohort MSE  $\downarrow$  on HEST-1k, per cohort. Mean<sub>std</sub> over three seeds with nested-validation; best (lowest) per row in bold. Same models, cohorts, features (frozen UNI + CONCH, 1536-dim), 50-gene panels, and training budget as Table 2.

| Cohort          | ST-Net NBE'20     | Hist2ST BIB'22                    | BLEEP NeurIPS'23                  | STEM ICLR'25      | MIST (Ours)                       |
|-----------------|-------------------|-----------------------------------|-----------------------------------|-------------------|-----------------------------------|
| IDC Breast      | 2.966 $\pm$ 0.311 | 2.740 $\pm$ 0.194                 | 2.503 $\pm$ 0.061                 | 2.811 $\pm$ 0.056 | <b>2.242<math>\pm</math>0.031</b> |
| COAD Colorectal | 2.335 $\pm$ 0.063 | <b>2.195<math>\pm</math>0.058</b> | 2.201 $\pm$ 0.062                 | 2.632 $\pm$ 0.021 | 2.322 $\pm$ 0.031                 |
| CCRCC Kidney    | 0.587 $\pm$ 0.007 | 0.578 $\pm$ 0.032                 | 0.604 $\pm$ 0.011                 | 0.763 $\pm$ 0.016 | <b>0.556<math>\pm</math>0.019</b> |
| LUNG Lung       | 3.388 $\pm$ 0.062 | 1.732 $\pm$ 0.003                 | <b>1.650<math>\pm</math>0.024</b> | 2.036 $\pm$ 0.052 | 1.665 $\pm$ 0.043                 |
| PAAD Pancreas   | 2.051 $\pm$ 0.235 | 1.331 $\pm$ 0.193                 | <b>1.151<math>\pm</math>0.258</b> | 1.711 $\pm$ 0.206 | 1.173 $\pm$ 0.166                 |
| PRAD Prostate   | 0.586 $\pm$ 0.011 | 0.560 $\pm$ 0.020                 | 0.570 $\pm$ 0.019                 | 0.822 $\pm$ 0.014 | <b>0.537<math>\pm</math>0.034</b> |
| SKCM Skin       | 2.387 $\pm$ 0.085 | 2.176 $\pm$ 0.016                 | 1.570 $\pm$ 0.036                 | 2.207 $\pm$ 0.260 | <b>1.227<math>\pm</math>0.030</b> |
| Average         | 2.043             | 1.616                             | 1.464                             | 1.855             | <b>1.389</b>                      |

Table 21: Intra-cohort MAE  $\downarrow$  on HEST-1k, per cohort. Mean<sub>std</sub> over three seeds with nested-validation; best (lowest) per row in bold. Same models, cohorts, features (frozen UNI + CONCH, 1536-dim), 50-gene panels, and training budget as Table 2.

| Cohort          | ST-Net NBE'20     | Hist2ST BIB'22                    | BLEEP NeurIPS'23                  | STEM ICLR'25      | MIST+CONCH                        |
|-----------------|-------------------|-----------------------------------|-----------------------------------|-------------------|-----------------------------------|
| IDC Breast      | 1.273 $\pm$ 0.066 | 1.246 $\pm$ 0.032                 | 1.141 $\pm$ 0.022                 | 1.210 $\pm$ 0.015 | <b>1.096<math>\pm</math>0.010</b> |
| COAD Colorectal | 0.996 $\pm$ 0.014 | 1.008 $\pm$ 0.017                 | <b>0.986<math>\pm</math>0.020</b> | 1.088 $\pm$ 0.008 | 1.040 $\pm$ 0.012                 |
| CCRCC Kidney    | 0.575 $\pm$ 0.003 | 0.574 $\pm$ 0.018                 | 0.583 $\pm$ 0.003                 | 0.612 $\pm$ 0.011 | <b>0.569<math>\pm</math>0.008</b> |
| LUNG Lung       | 1.373 $\pm$ 0.014 | 0.985 $\pm$ 0.002                 | 0.960 $\pm$ 0.008                 | 1.050 $\pm$ 0.003 | <b>0.957<math>\pm</math>0.014</b> |
| PAAD Pancreas   | 1.013 $\pm$ 0.051 | 0.872 $\pm$ 0.055                 | <b>0.796<math>\pm</math>0.088</b> | 0.936 $\pm$ 0.054 | 0.798 $\pm$ 0.046                 |
| PRAD Prostate   | 0.564 $\pm$ 0.004 | <b>0.546<math>\pm</math>0.003</b> | 0.547 $\pm$ 0.007                 | 0.620 $\pm$ 0.014 | 0.550 $\pm$ 0.017                 |
| SKCM Skin       | 1.015 $\pm$ 0.023 | 1.056 $\pm$ 0.015                 | 0.926 $\pm$ 0.011                 | 1.049 $\pm$ 0.070 | <b>0.791<math>\pm</math>0.005</b> |
| Average         | 0.973             | 0.898                             | 0.848                             | 0.938             | <b>0.829</b>                      |

## L Cross-Organ Generalization: PCC and MSE

Table 22 reports both PCC and MSE for the LOOO and POOLED cross-organ regimes, complementing the PCC-only summaries in the main paper.

Table 22: Cross-organ generalization: PCC and MSE (three seeds, mean<sub>std</sub>). looo: leave-one-organ-out transfer; pooled: 5-fold pooled cross-validation. PCC $\uparrow$  / MSE $\downarrow$  on HEST-1k.

| Method                               | looo                        |                             | pooled                      |                             |
|--------------------------------------|-----------------------------|-----------------------------|-----------------------------|-----------------------------|
|                                      | PCC $\uparrow$              | MSE $\downarrow$            | PCC $\uparrow$              | MSE $\downarrow$            |
| ST-Net He et al. (2020) (NBE'20)     | 0.267 <sub>0.02</sub>       | 1.846 <sub>0.10</sub>       | 0.370 <sub>0.01</sub>       | 0.907 <sub>0.05</sub>       |
| Hist2ST Zeng et al. (2022) (BIB'22)  | 0.360 <sub>0.02</sub>       | 1.182 <sub>0.07</sub>       | 0.737 <sub>0.01</sub>       | 0.488 <sub>0.01</sub>       |
| BLEEP Xie et al. (2023) (NeurIPS'23) | 0.416 <sub>0.01</sub>       | <b>1.010<sub>0.02</sub></b> | 0.758 <sub>0.01</sub>       | <b>0.427<sub>0.02</sub></b> |
| STEM Zhu et al. (2025) (ICLR'25)     | 0.322 <sub>0.01</sub>       | 1.274 <sub>0.00</sub>       | 0.643 <sub>0.01</sub>       | 0.594 <sub>0.01</sub>       |
| MIST (ours)                          | <b>0.421<sub>0.01</sub></b> | 1.233 <sub>0.11</sub>       | <b>0.760<sub>0.01</sub></b> | 0.440 <sub>0.02</sub>       |

## M Intra-Cohort Results with UNI + CONCH Features, Extended Baselines

Table 23 extends the intra-cohort comparison to include M2OST and EGN under the same UNI + CONCH feature protocol.

Table 23: Intra-cohort results using UNI+CONCH features, extended baselines (Pearson  $\text{mean}_{\text{std}}$  over three seeds; nested-validation). All methods use identical frozen UNI+CONCH features (1536-dim) and 50-gene panels. HCC (liver), LYMPH\_IDC, and READ are excluded per Appendix H. Best per row in bold.

| Cohort          | ST-Net     | NBE'20     | Hist2ST    | BIB'22            | BLEEP      | NeurIPS'23 | M2OST             | AAAI'25 | EGN | WACV'23 | STEM | ICLR'25 | MIST (ours) |
|-----------------|------------|------------|------------|-------------------|------------|------------|-------------------|---------|-----|---------|------|---------|-------------|
| IDC Breast      | 0.551±0.01 | 0.309±0.06 | 0.519±0.00 | 0.571±0.01        | 0.551±0.01 | 0.439±0.01 | <b>0.581±0.01</b> |         |     |         |      |         |             |
| COAD Colorectal | 0.242±0.01 | 0.295±0.01 | 0.314±0.01 | <b>0.328±0.01</b> | 0.293±0.01 | 0.177±0.01 | 0.318±0.01        |         |     |         |      |         |             |
| CCRCC Kidney    | 0.163±0.02 | 0.184±0.02 | 0.175±0.02 | 0.210±0.02        | 0.223±0.02 | 0.123±0.00 | <b>0.252±0.01</b> |         |     |         |      |         |             |
| LUNG Lung       | 0.518±0.01 | 0.470±0.01 | 0.562±0.01 | <b>0.598±0.00</b> | 0.584±0.01 | 0.423±0.03 | 0.589±0.01        |         |     |         |      |         |             |
| PAAD Pancreas   | 0.359±0.03 | 0.247±0.02 | 0.419±0.03 | 0.413±0.04        | 0.415±0.03 | 0.233±0.02 | <b>0.437±0.03</b> |         |     |         |      |         |             |
| PRAD Prostate   | 0.294±0.02 | 0.333±0.02 | 0.348±0.01 | <b>0.413±0.01</b> | 0.377±0.03 | 0.207±0.02 | 0.410±0.01        |         |     |         |      |         |             |
| SKCM Skin       | 0.665±0.01 | 0.522±0.03 | 0.645±0.01 | 0.654±0.01        | 0.631±0.01 | 0.491±0.04 | <b>0.707±0.01</b> |         |     |         |      |         |             |
| Average         | 0.399      | 0.337      | 0.426      | 0.455             | 0.439      | 0.299      | <b>0.470</b>      |         |     |         |      |         |             |

All entries use three completed seeds. STEM uses the corrected DiT configuration ( $T = 1000$ , 100 epochs).

## N Training Objective Ablation

Table 24 isolates the contribution of the Pearson correlation term in the training objective from the architectural design, confirming that the accuracy gap over baselines is not attributable to the loss function alone.

Table 24: Objective ablation (HEST-1k). Effect of adding the Pearson correlation term to the MSE objective; averaged over seven organ groups, three seeds.

| Model | Objective | PCC $\uparrow$ |
|-------|-----------|----------------|
| MIST  | MSE only  | 0.382          |
| MIST  | MSE + PCC | 0.389          |